The background is a solid blue color with a faint, abstract network diagram. It consists of numerous small blue dots (nodes) connected by thin, light blue lines (edges). The nodes are arranged in a way that suggests a complex, interconnected system, with some nodes having many connections and others having fewer. The overall pattern is organic and radial, resembling a star or a web.

Slides, Support, Code and Setup instructions:
bit.ly/hub-setup



Introductions

Keras/LSTM Text Classification Agenda

9:00 - 9:45

Keras/CNN Review

examples/keras-sign

9:45 - 10:30

Simple RNNs and Time Series

examples/lstm/time-series

10:30 - 11:00

Break

10:30 - 11:00

LSTMs, GRUs applied to Text Generation

examples/lstm/text-gen

11:00 - 12:30

Text Classification with Word Embeddings, Hybrid
CNN/LSTMs

examples/lstm/imdb-classifier

Code: github.com/lukas/ml-class



Keras

TensorFlow

CudNN

CUDA


GPU

Build your own box

Build a super fast deep learning machine for under \$1,000

The adventures in deep learning and cheap hardware continue!

By Lukas Biewald. February 1, 2017




<https://www.oreilly.com/learning/build-a-super-fast-deep-learning-machine-for-under-1000>



Keras Review ([ml-class/examples/keras-sign/](https://ml-class.github.io/examples/keras-sign/))

Perceptron



Schematic of Rosenblatt's perceptron.

perceptron-single.py

0	91	95	0
0	0	99	0
0	0	35	0
0	0	94	0



Input (28x28)



Flatten (1x784)

0	91	95	0	0	0	99	0	0	0	35	0	0	0	94	0
---	----	----	---	---	---	----	---	---	---	----	---	---	---	----	---



Σ

0.3

1.0

Dense (1x1)

Target

perceptron-linear.py

0	91	95	0
0	0	99	0
0	0	35	0
0	0	94	0




Input (28x28)



Flatten (1x784)

0	91	95	0	0	0	99	0	0	0	35	0	0	0	94	0
---	----	----	---	---	---	----	---	---	---	----	---	---	---	----	---



Dense (1x10)

0	1	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---


Target

One hot encoding


Label
0
4
4
3
0
9


[illegible]

Activation Functions: Sigmoid




(special case of softmax and logistic)









Two Layers of Perceptrons





MLP Activation Function: ReLU




Overfitting








Pooling









Time Series and RNNs

([ml-class/examples/lstm/time-series](#))


Timeseries (Airline Sales)



Timeseries (Airline Sales)



Timeseries (Airline Sales)



Framing Timeseries as an ML problem

Time	1	2	3	4	5	6	7	8	9	10	11	12	13
Airline Sales	1	3	4	7	11	18	29	31	42	55	62	74	78

Framing Timeseries as an ML problem

Time	1	2	3	4	5	6	7	8	9	10	11	12	13
Airline Sales	1	3	4	7	11	18	29	31	42	55	62	74	78

Train

Label



X1	X2	X3-9	X10	Label
1	3	...	55	62

Framing Timeseries as an ML problem

Time	1	2	3	4	5	6	7	8	9	10	11	12	13
Airline Sales	1	3	4	7	11	18	29	31	42	55	62	74	78

Train

Label



X1	X2	X3-9	X10	Label
1	3	...	55	62
3	4	...	62	74

Framing Timeseries as an ML problem

Time	1	2	3	4	5	6	7	8	9	10	11	12	13
Airline Sales	1	3	4	7	11	18	29	31	42	55	62	74	78

Train


Label




X1	X2	X3-9	X10	Label
1	3	...	55	62
3	4	...	62	74
4	7	...	74	78

Perceptron Classifier on Timeseries

X1	X2	X3-9	X10	Label
1	3	...	55	62
3	4	...	62	74
4	7	...	74	78



Input




Dense (1x1)

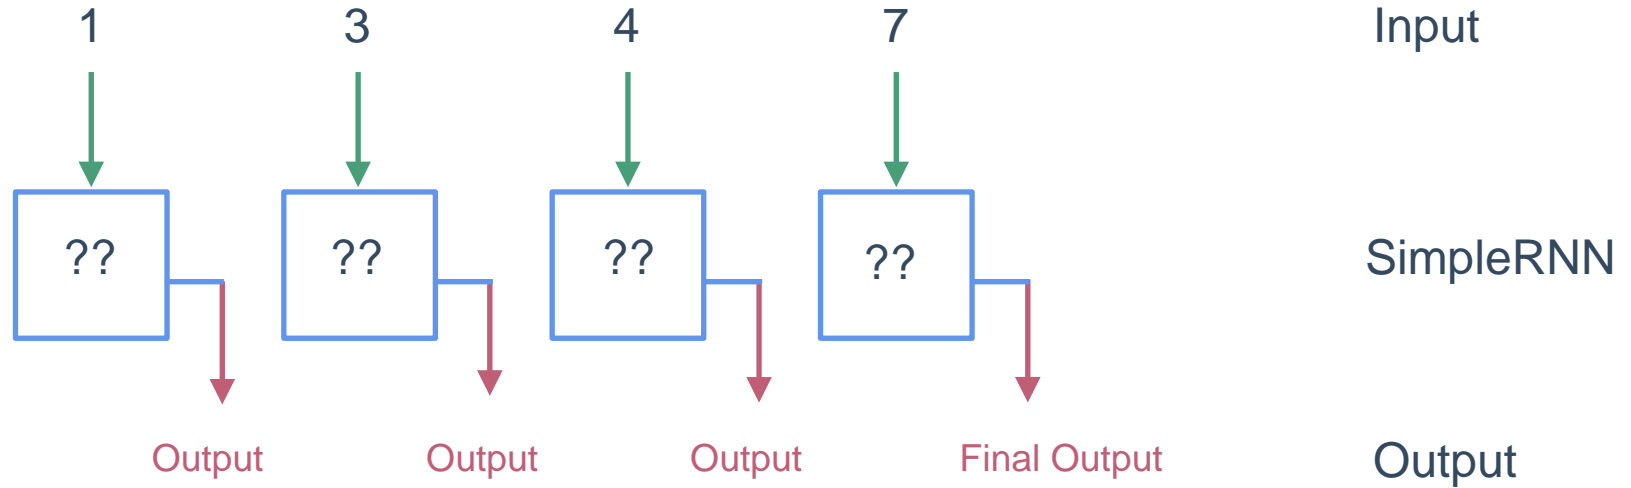


Target

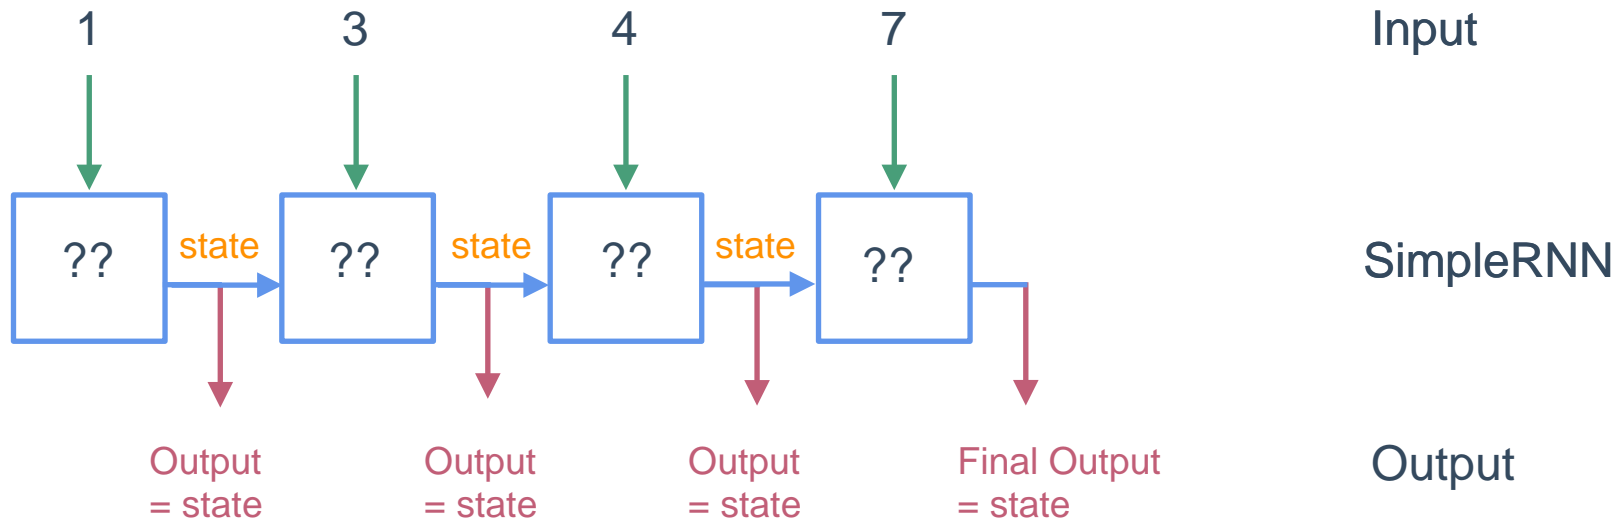
Recurrent Neural Network (RNN)




Recurrent Neural Network (RNN)



Recurrent Neural Network (RNN)



RNN



Input

SimpleRNN

Output

Inside the RNN

State	Input
0.4	3


Input

0.2	0.4
-----	-----

Learned Weights

$$0.2 * 0.4 + 0.4 * 3 = 1.28$$

Weighted Sum




Activation Function

0.9

Output

RNN




Input

SimpleRNN

Output

Recurrent Neural Network (RNN)




Input

SimpleRNN

Output

Recurrent Neural Network (RNN)




Input

SimpleRNN

Output

Recurrent Neural Network (RNN)

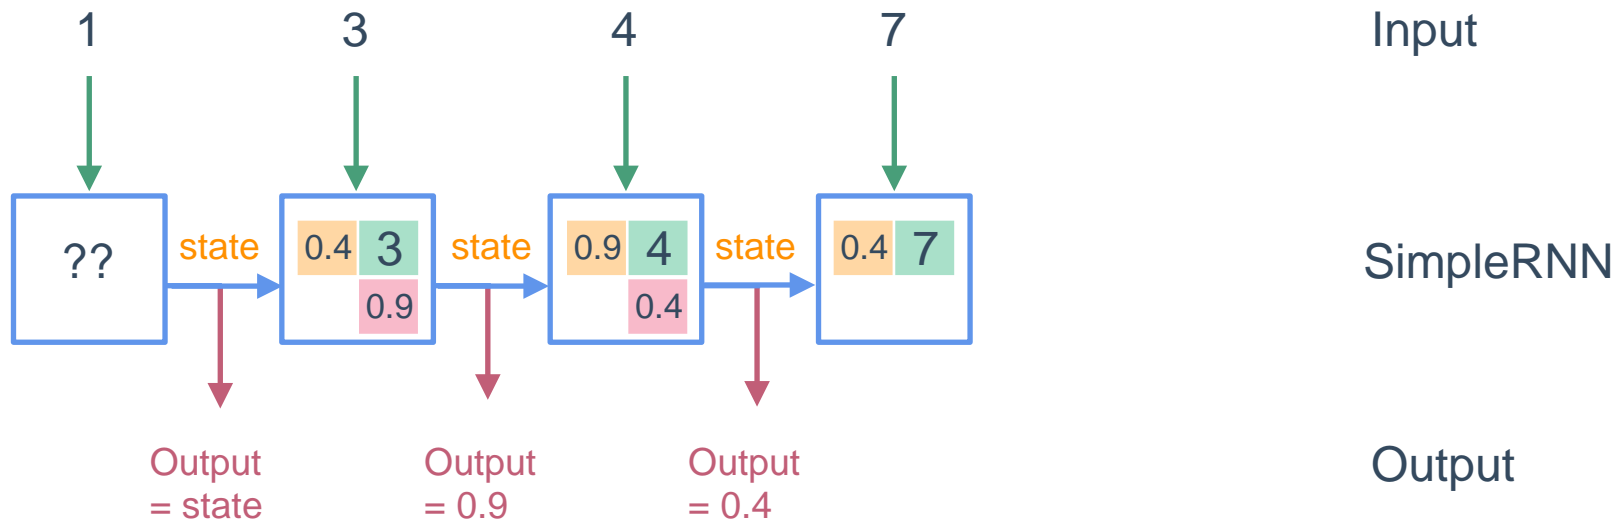


Input

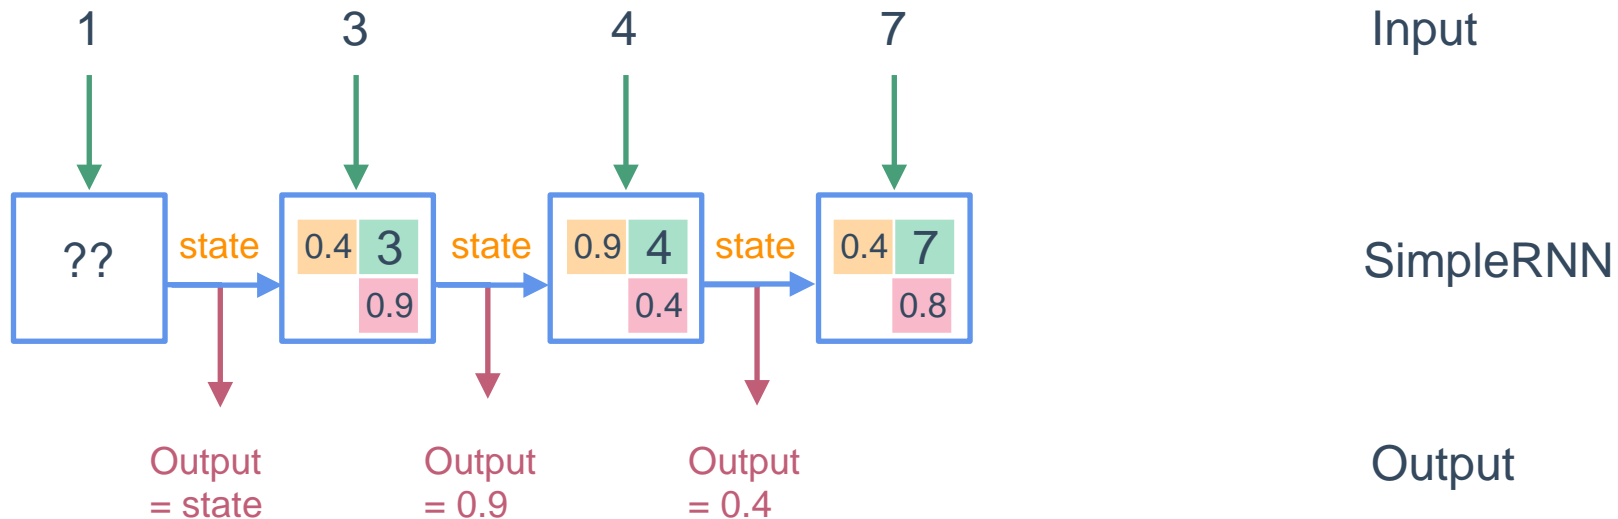
SimpleRNN

Output

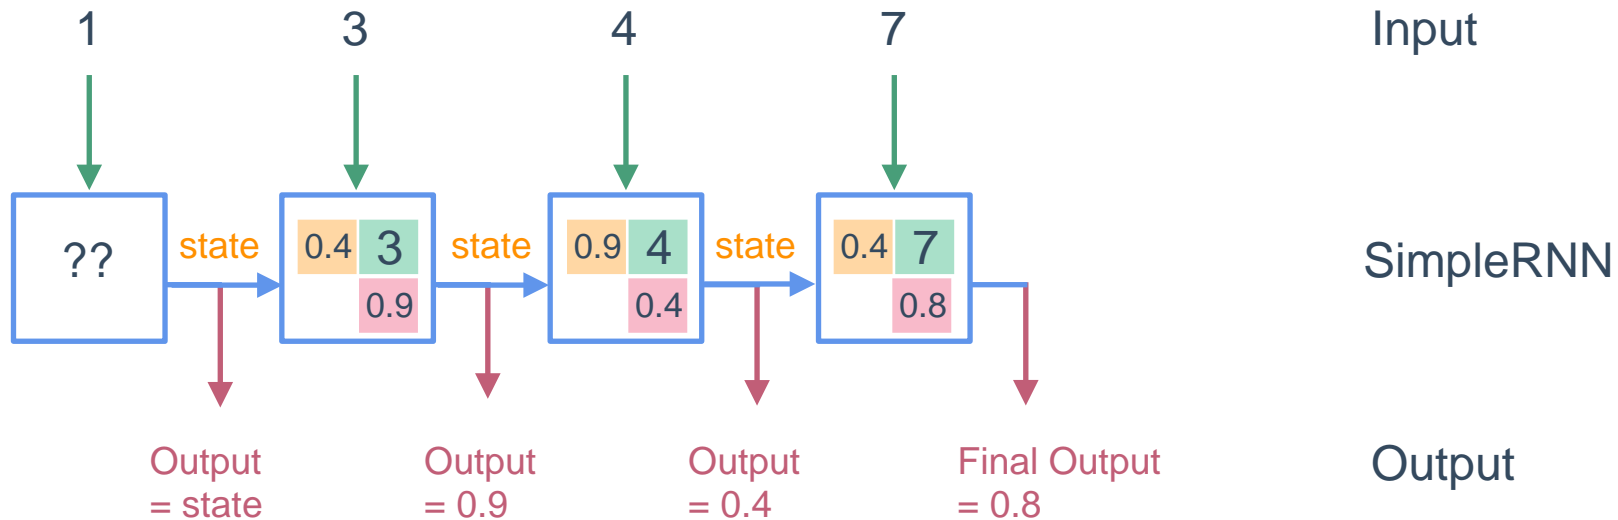
Recurrent Neural Network (RNN)




Recurrent Neural Network (RNN)




Recurrent Neural Network (RNN)



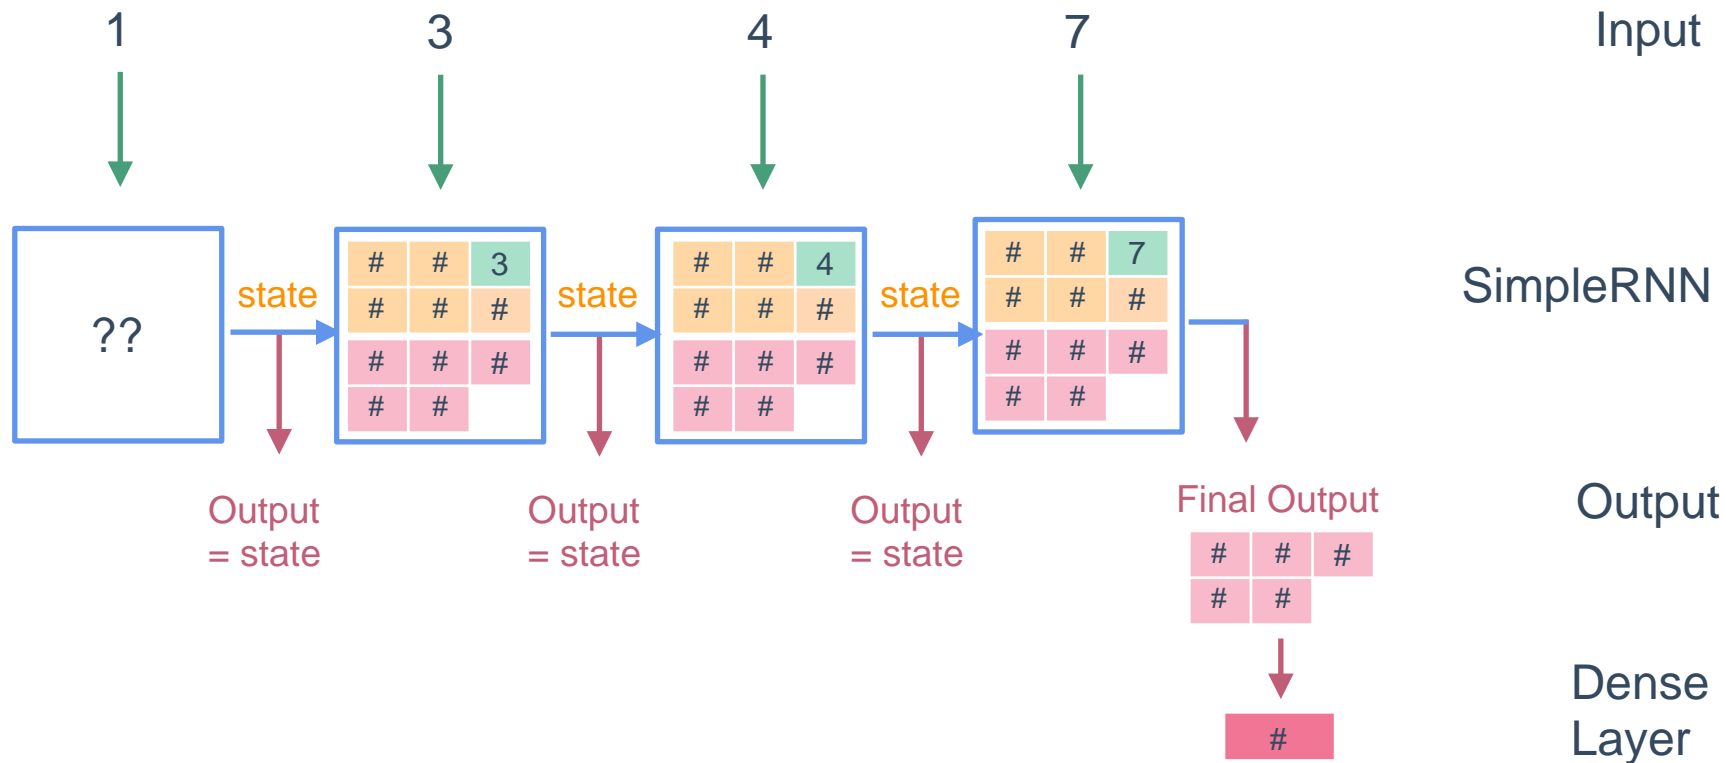
Hyperbolic Tangent Activation Function (tanh)



Hyperbolic Tangent Activation Function (tanh)



Multidimensional State RNN





LSTMs, GRUs & Text Generation

[ml-class/examples/lstm/text-gen](#)

Machine learning generated pranks

Put a pair of pants and shoes into your ice dispenser

Put marbles in the refrigerator

A meat and mash potato sundae makes for quite the hand soap dispenser

(Generated with an LSTM at aiweirdness.com)



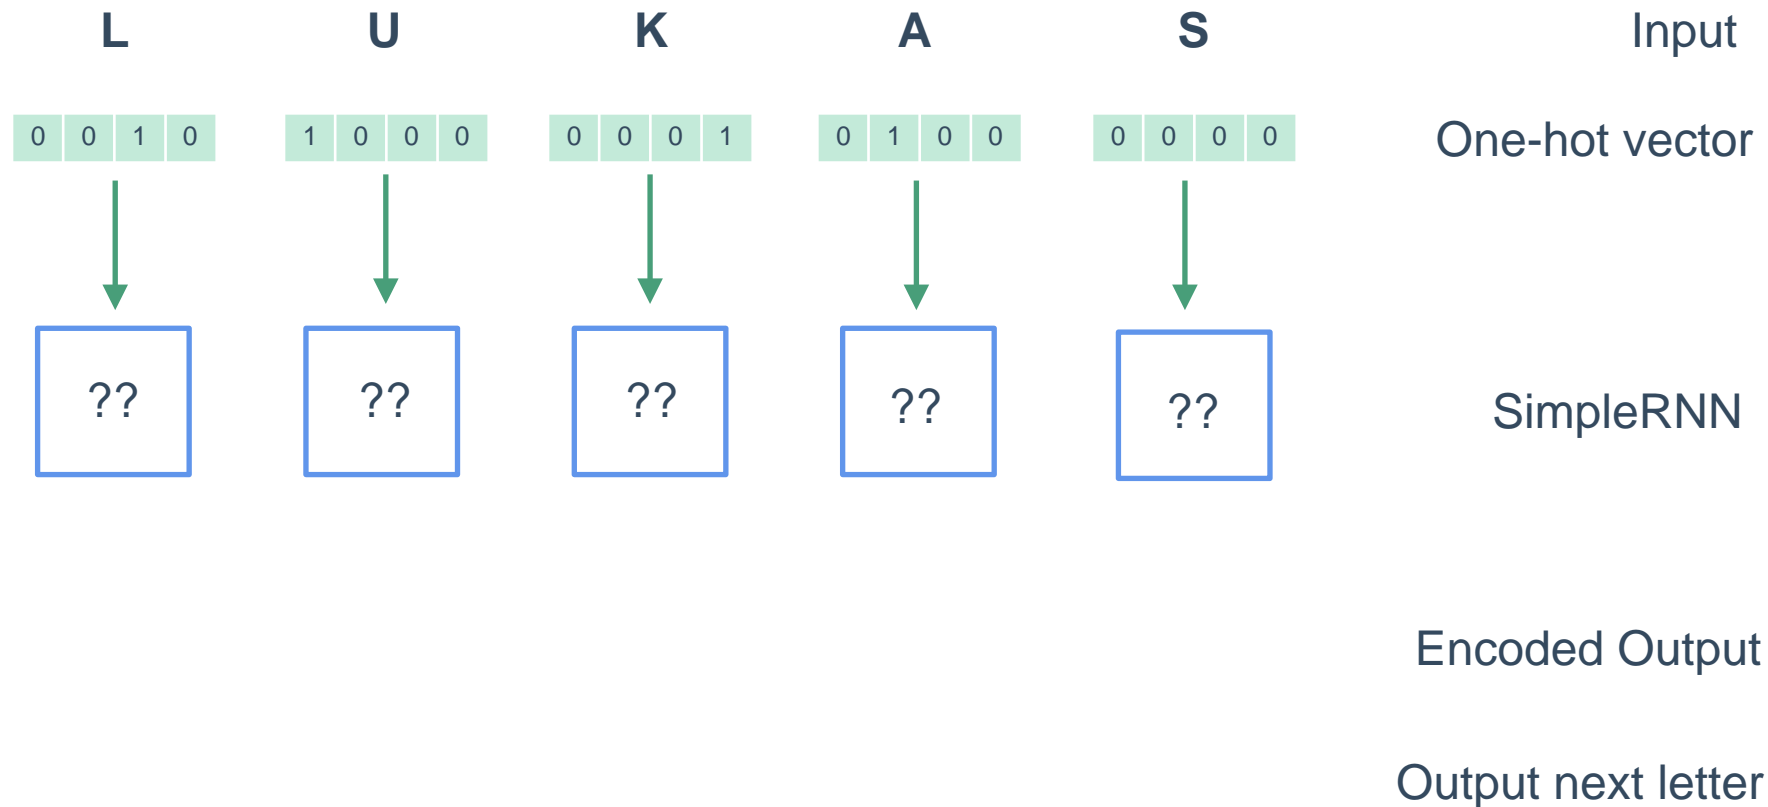
One-hot encoding

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
L →	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
U →	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
K →	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
A →	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
S →	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0

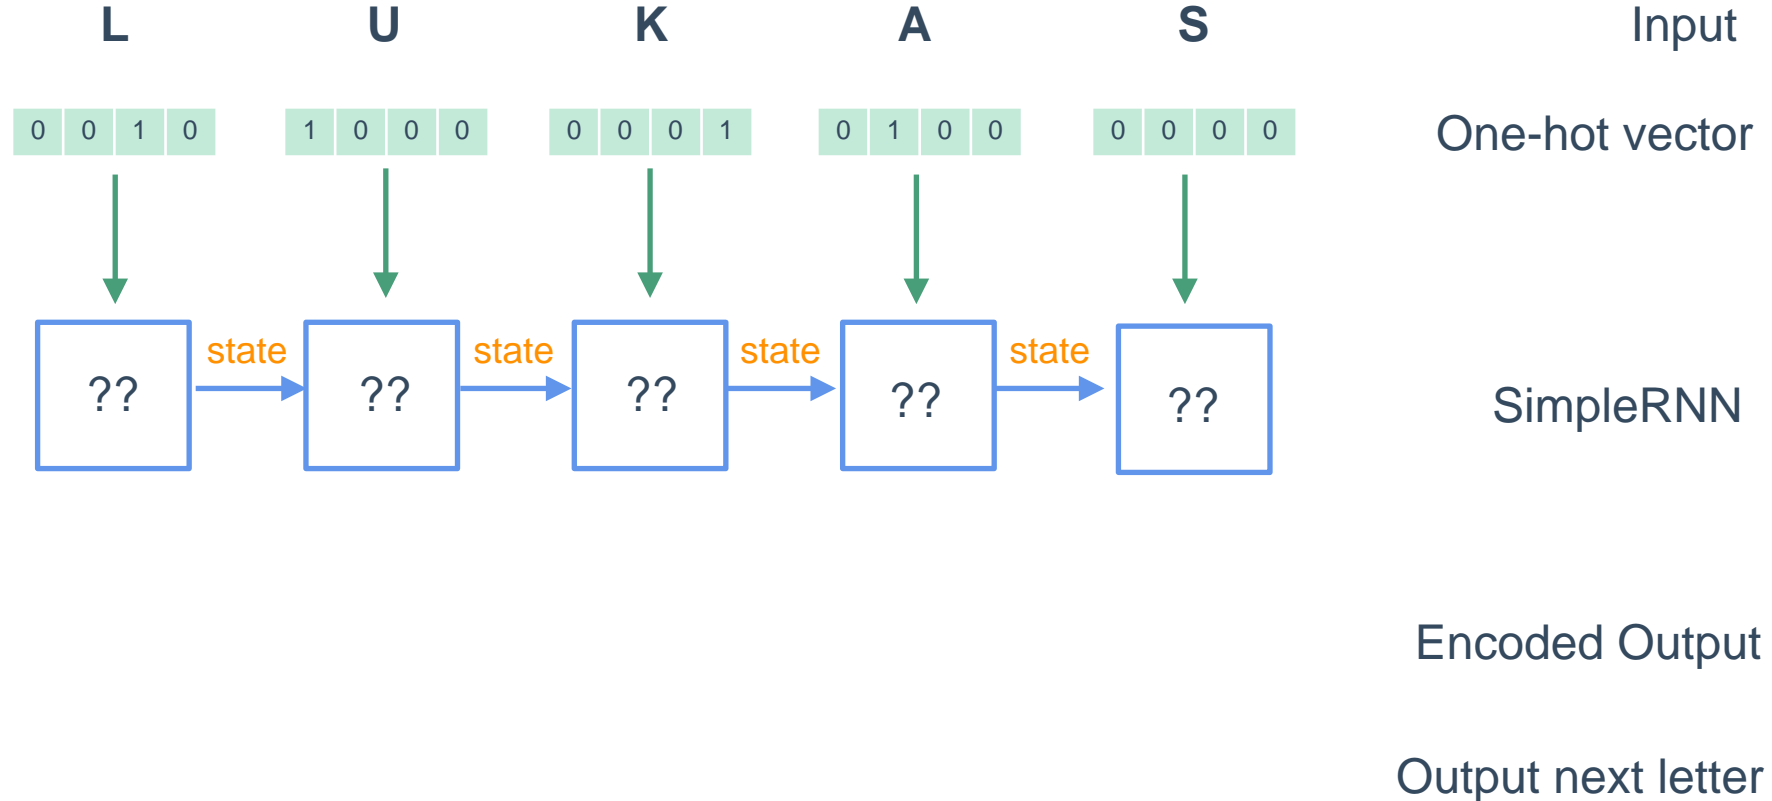
Text as time series (character-level)



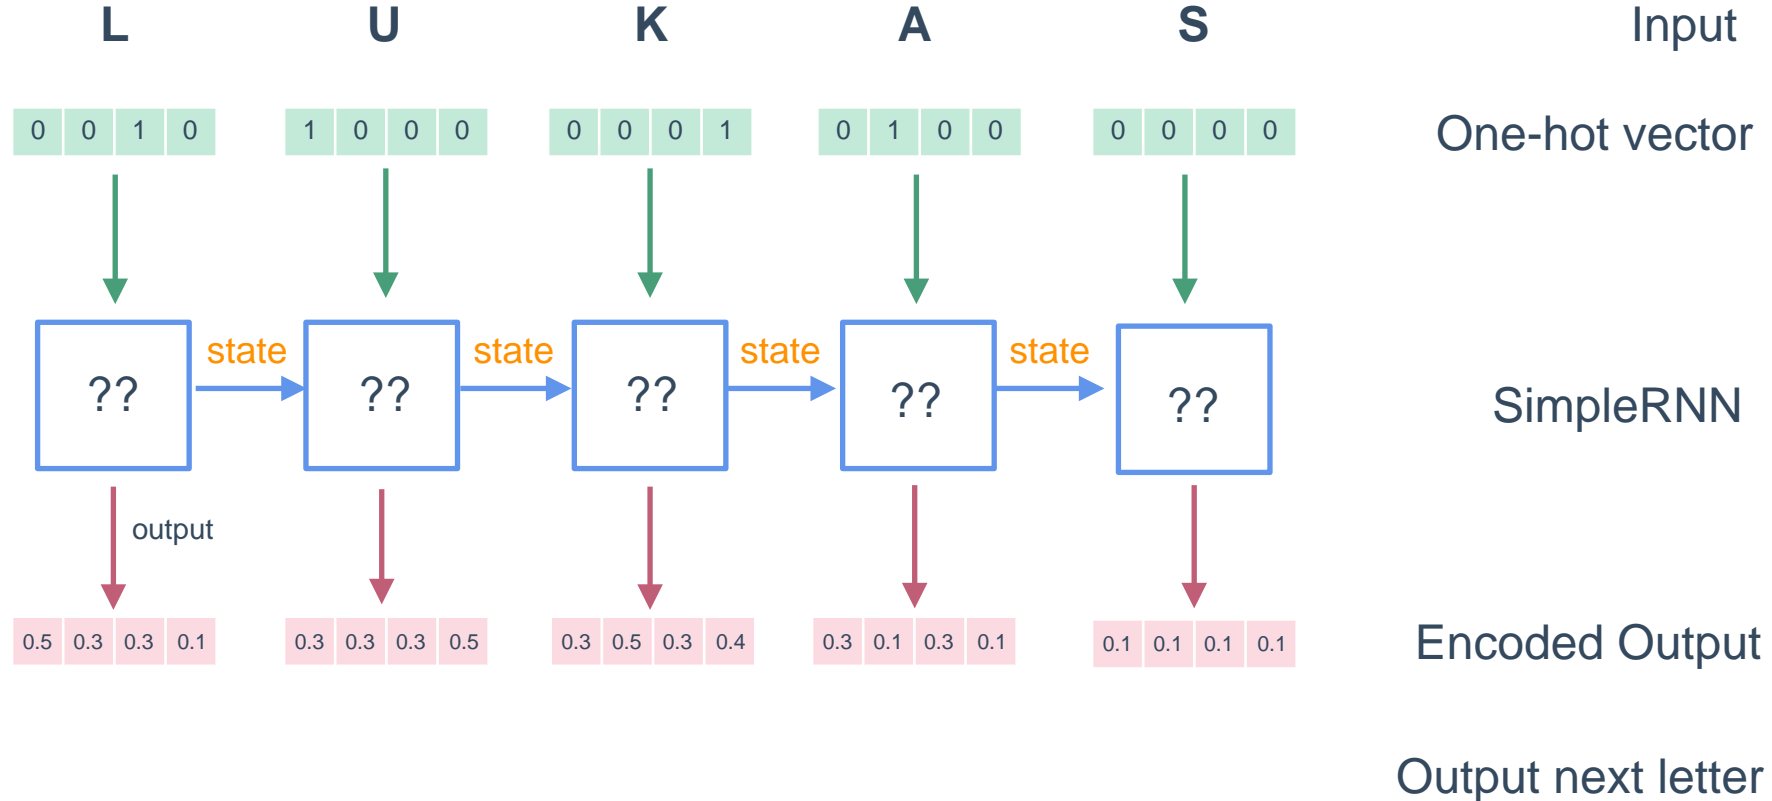
Text as time series (character-level)



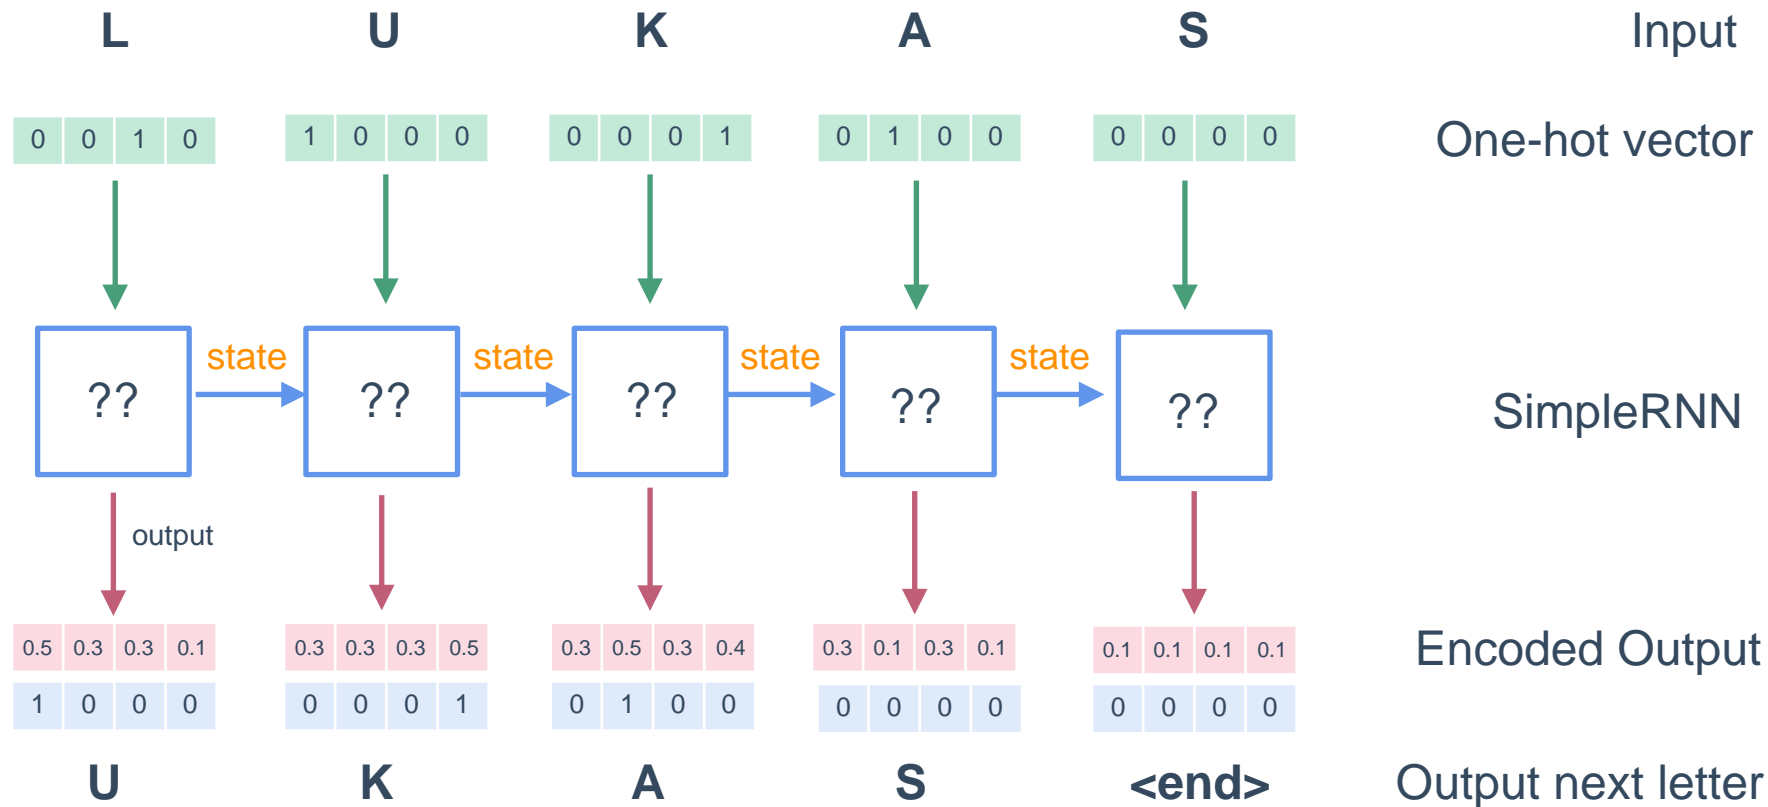
Text as time series (character-level)




Text as time series (character-level)



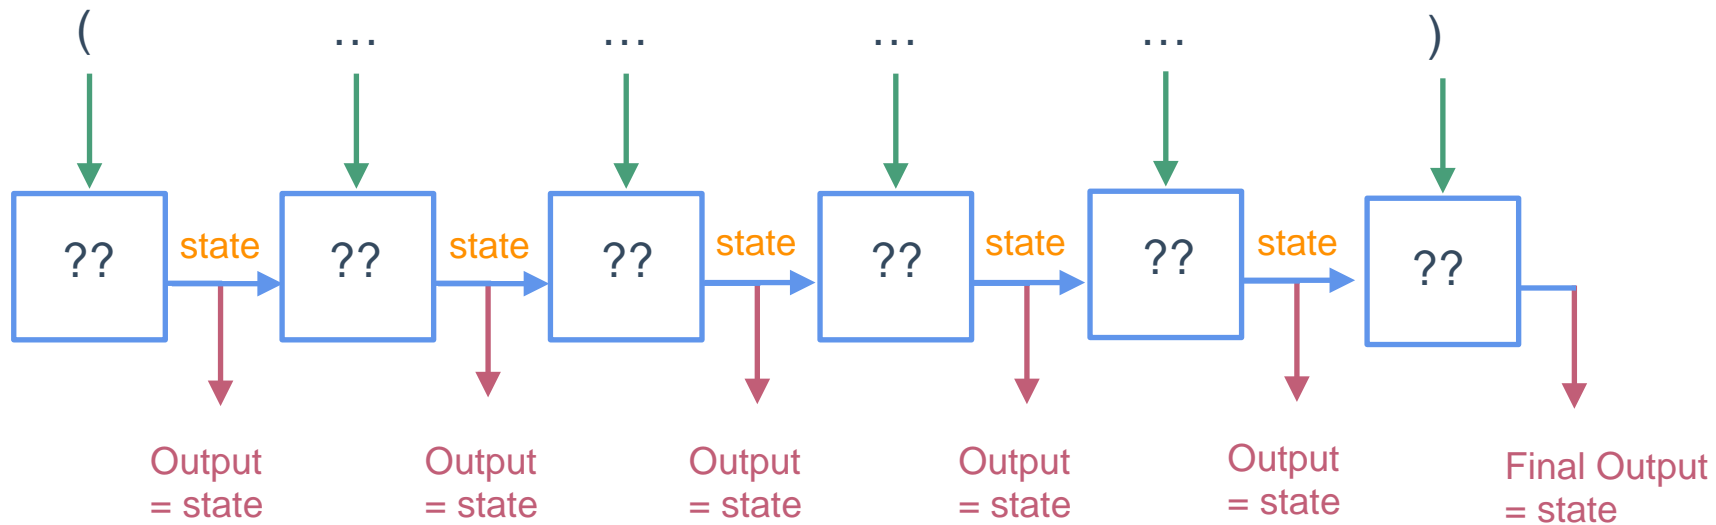
Text as time series (character-level)



The problem with simple RNNs



The problem with simple RNNs



Long Short Term Memory (LSTM) 1997

$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$


$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$


$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

$$h_t = o_t \circ \sigma_h(c_t)$$

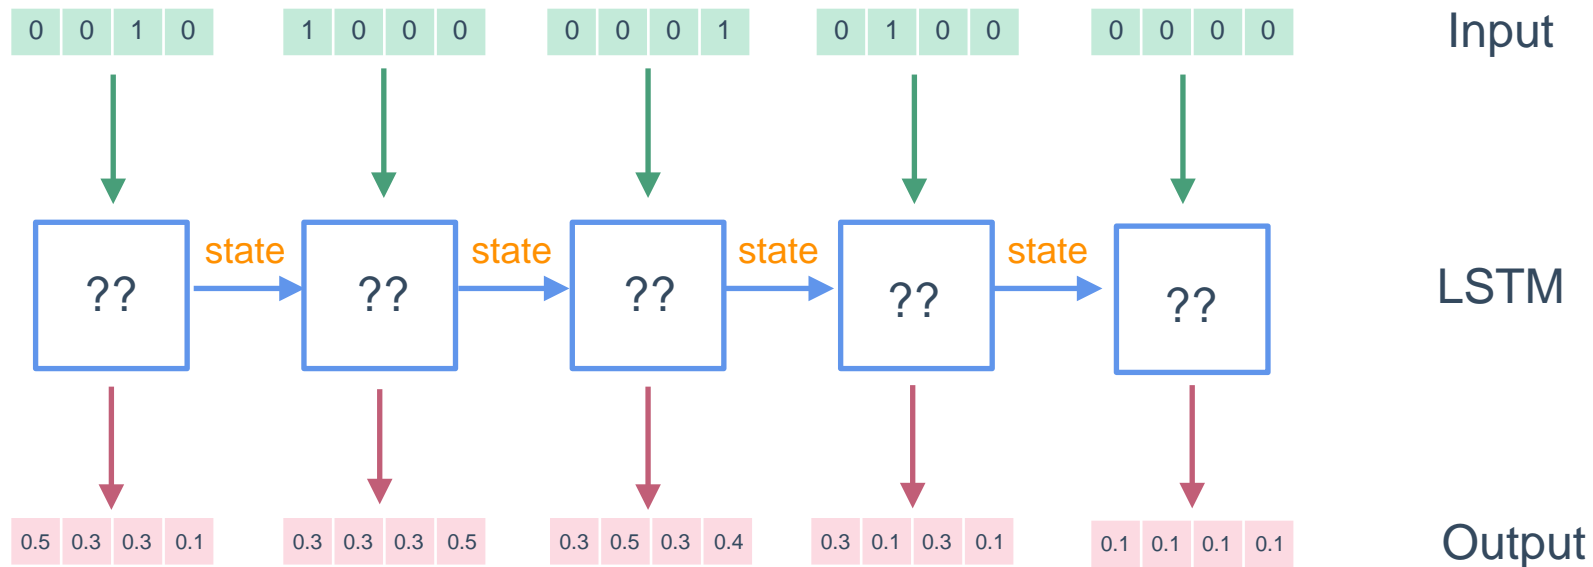
LSTM



LSTM



LSTM



LSTM

Input

0	0	1	0
---	---	---	---



LSTM

Input

0	0	1	0
---	---	---	---




Forgetting



LSTM

Input

0	0	1	0
---	---	---	---




Forgetting

Update Activation

LSTM

Input


0 0 1 0



LSTM

Input

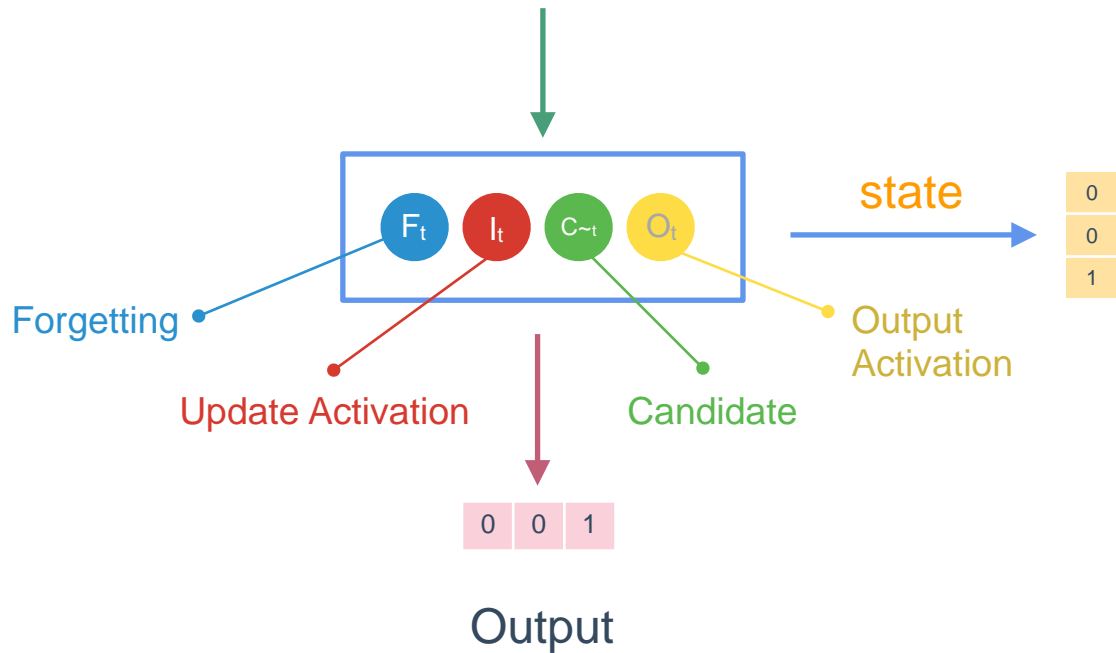
0 0 1 0



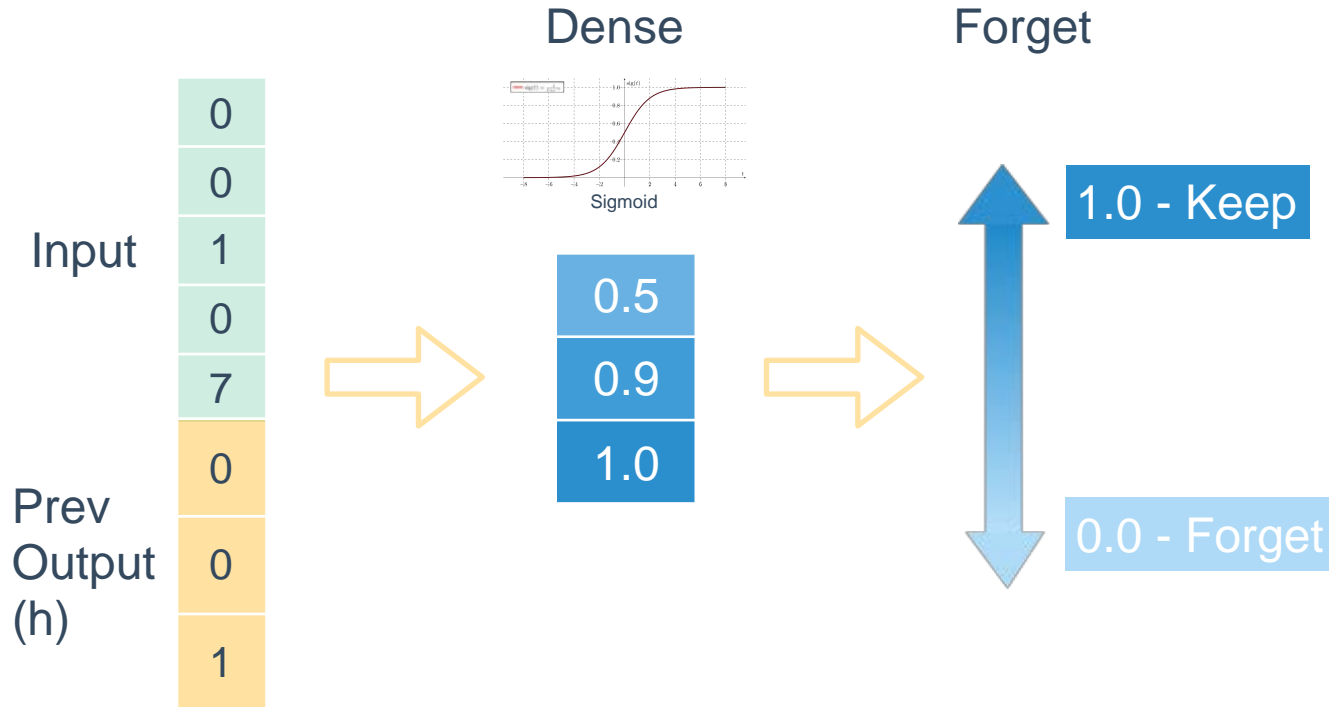
LSTM

Input

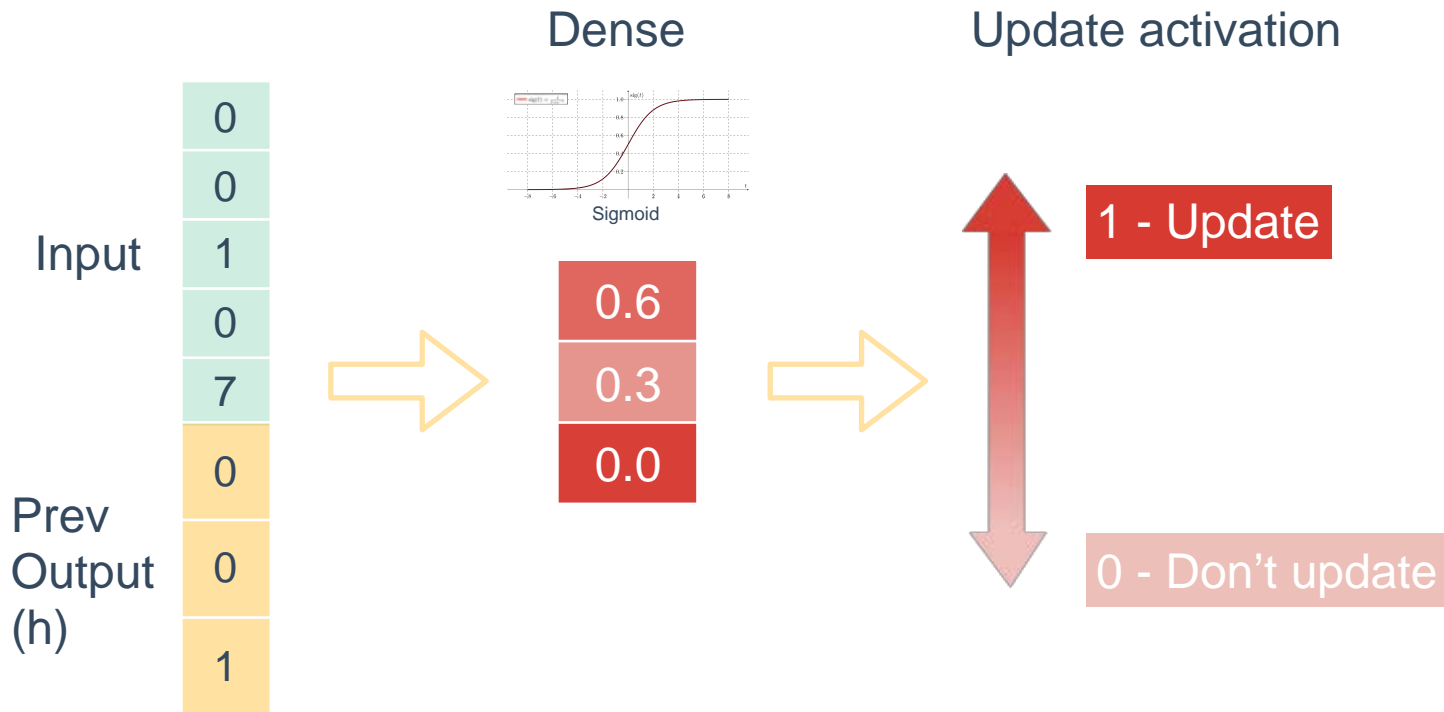
0 0 1 0



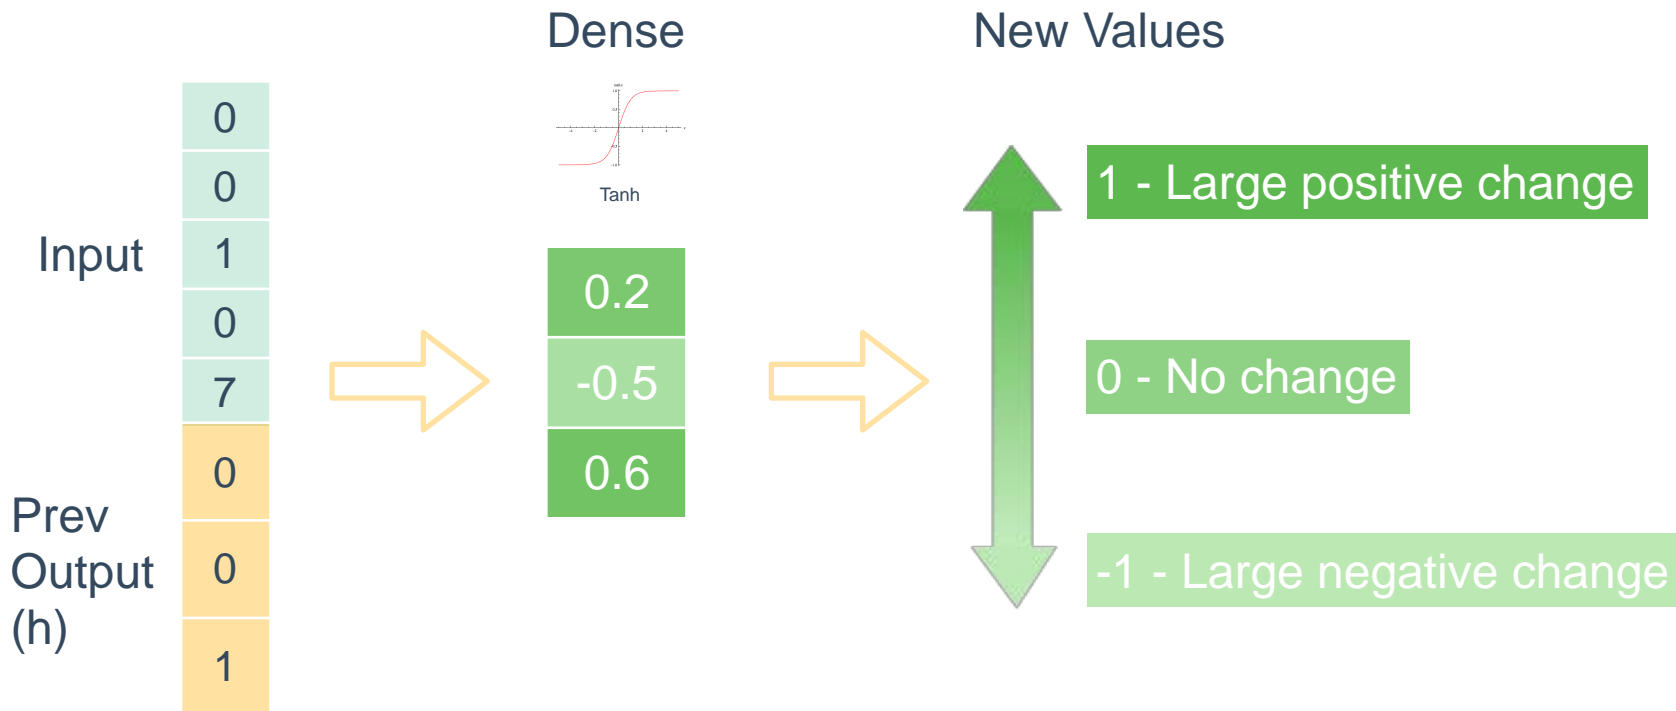
LSTM Forgetting (F_t)



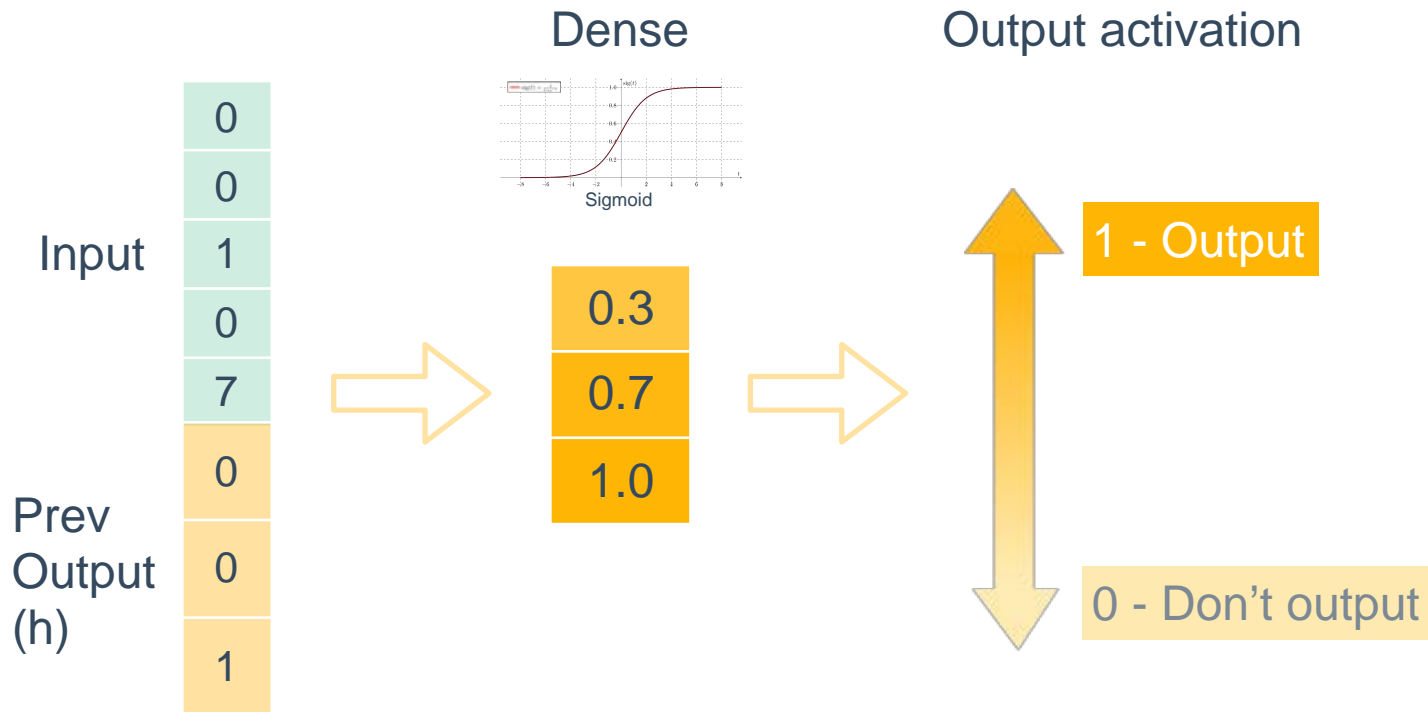
LSTM Update Activation (I_t)



LSTM Candidate (\tilde{C}_t)



LSTM Output Activation (O_t)





Putting it all together

Next State
(C_t)

??
??
??

=

Prev State
(C_{t-1})

0
0
1

*

Forget

F_t

0.5
0.9
1.0



Putting it all together

Next State
(C_t)

??
??
??

=

Prev State
(C_{t-1})

0
0
1

*

Forget

F_t

0.5
0.9
1.0

+

Update
Activation

I_t

0.6
0.3
0.0



*

Candidate


C_{-t}

0.2
-0.5
0.6

Putting it all together



Hard Sigmoid



Gated Recurrent Unit (GRU) 2014


$$z_t = \sigma(x_t U^z + h_{t-1} W^z)$$

$$r_t = \sigma(x_t U^r + h_{t-1} W^r)$$


$$\tilde{h}_t = \tanh(x_t U^h + (r_t * h_{t-1}) W^h)$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

LSTMs vs GRUs (1)




LSTM: State \neq Output




GRU: State $=$ Output

LSTMs vs GRUs (2)



LSTM: 4 vectors



GRU: 3 vectors



Text Classification, Embeddings and 1D Convolutions

[ml-class/examples/lstm/imdb-classifier](#)

Character encoding

		A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
L	→	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
U	→	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
K	→	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
A	→	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
S	→	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0

Bag of Words

Input Text

“Bag of Words”

I love my iPhone →

I hate my iPhone →

a	...	hate	I	iPhone	love	my	...	zoo
0	...	0	1	1	1	1	...	0
0	...	1	1	1	0	1	...	0

Word Embedding

I	love	this	movie
↓	↓	↓	↓
I	love	this	movie
0.4	9.7	-0.7	0.4
23.4	5.4	-5.2	-3.2
-0.6	-9.4	-2.1	-9.9
-2.4	-1.2	-0.9	-1.9

Embeddings

	Val 1	Val 2	Val 3	Val 4
a	0.1	-0.3	1.7	2.4
aardvark	-2.3	4.1	-5.2	3.1
...
<unknown>	0.3	0.9	0.8	0.2

Pre-computing encoding

GloVe: Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher, Christopher D. Manning

Introduction

GloVe is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

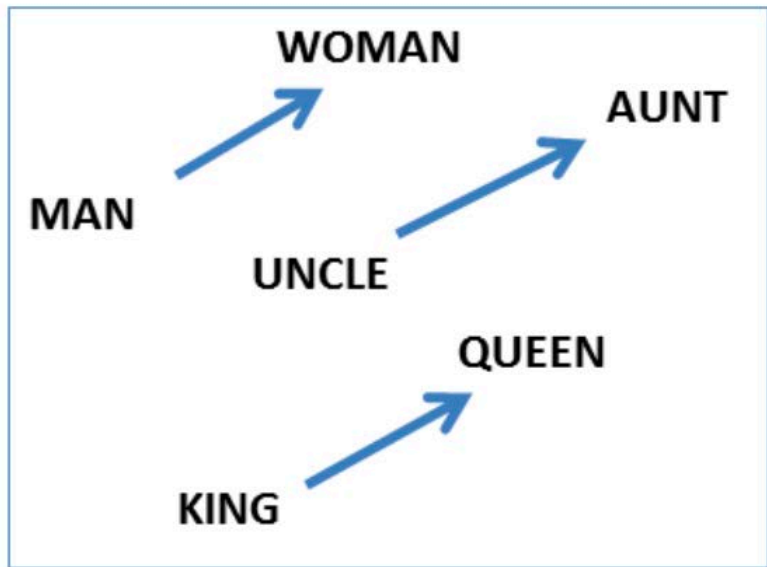
Getting started (Code download)

- Download the [code](#) (licensed under the [Apache License, Version 2.0](#))
- Unpack the files: `unzip GloVe-1.2.zip`
- Compile the source: `cd GloVe-1.2 && make`
- Run the demo script: `./demo.sh`
- Consult the included README for further usage details, or ask a [question](#)
- The code is also available [on GitHub](#)

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](#) v1.0 whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - [Wikipedia 2014](#) + [Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](#)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data

GloVe + word2vec



Word Analogy Task

man is to *woman* as *king* is to ____ ?

good is to *best* as *smart* is to ____ ?

china is to *beijing* as *russia* is to ____ ?

Turns out the word-context based vector model we just learnt is good for such analogy tasks,

$$[\text{king}] - [\text{man}] + [\text{woman}] \approx [\text{queen}]$$

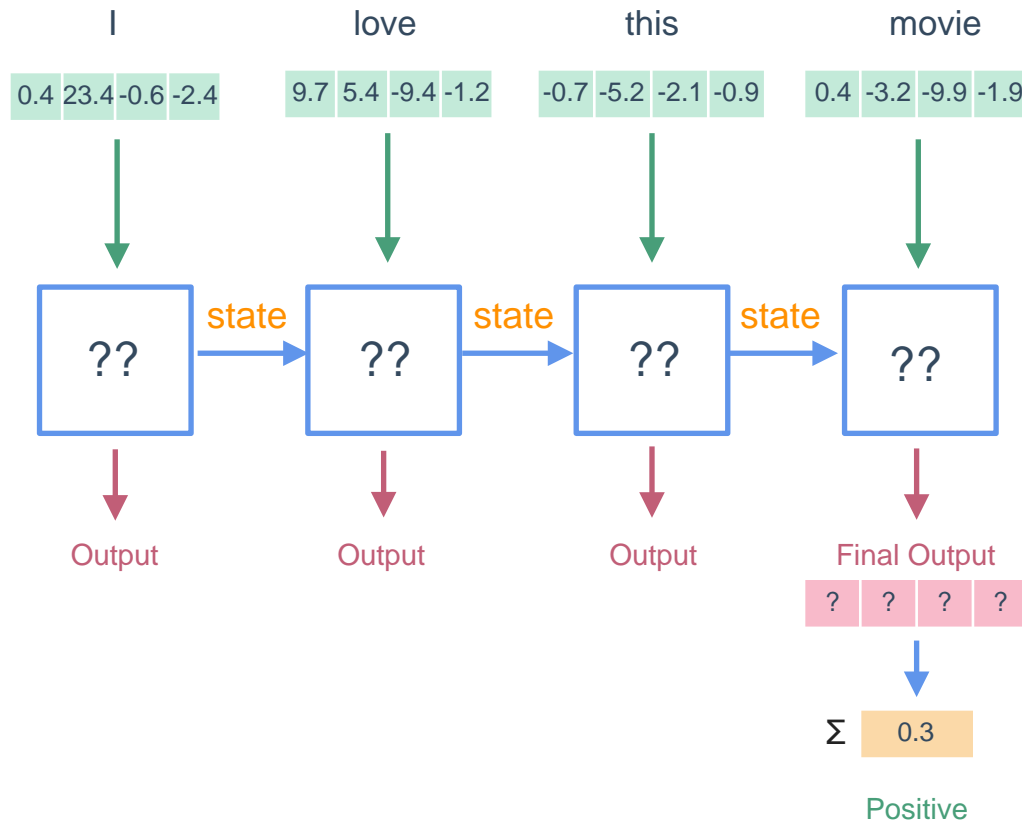


Levy, Goldberg, and Israel, [Linguistic Regularities in Sparse and Explicit Word Representations](#), CoNLL 2014.

The background is a solid blue color. Overlaid on this background is a faint, light-blue network diagram. This diagram consists of numerous small circular nodes connected by thin lines. The nodes are arranged in a way that suggests clusters or communities, with some nodes acting as hubs that connect to many other nodes, while others are more isolated or part of smaller groups. The overall pattern is complex and interconnected, resembling a social network or a data graph.

Classification with LSTMs

Classification LSTM



Input

Word Embeddings

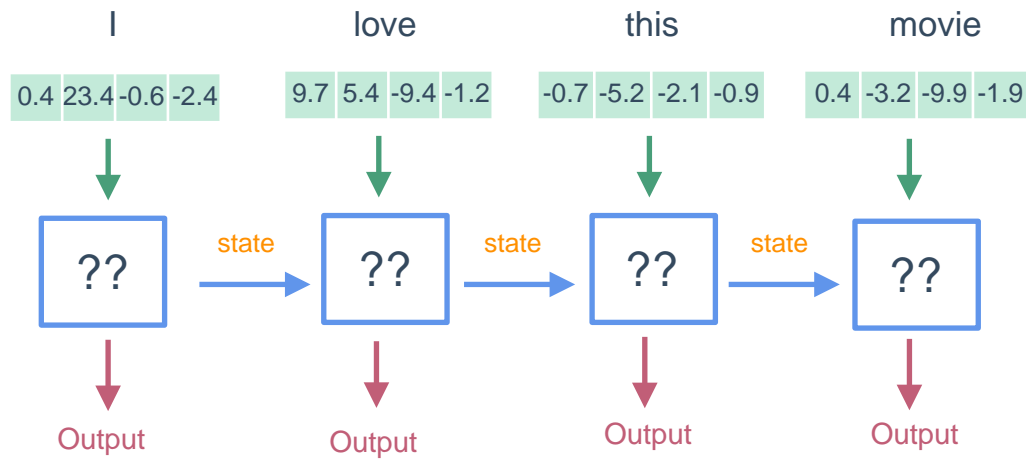
LSTM

Encoded Output

Dense layer

Classification

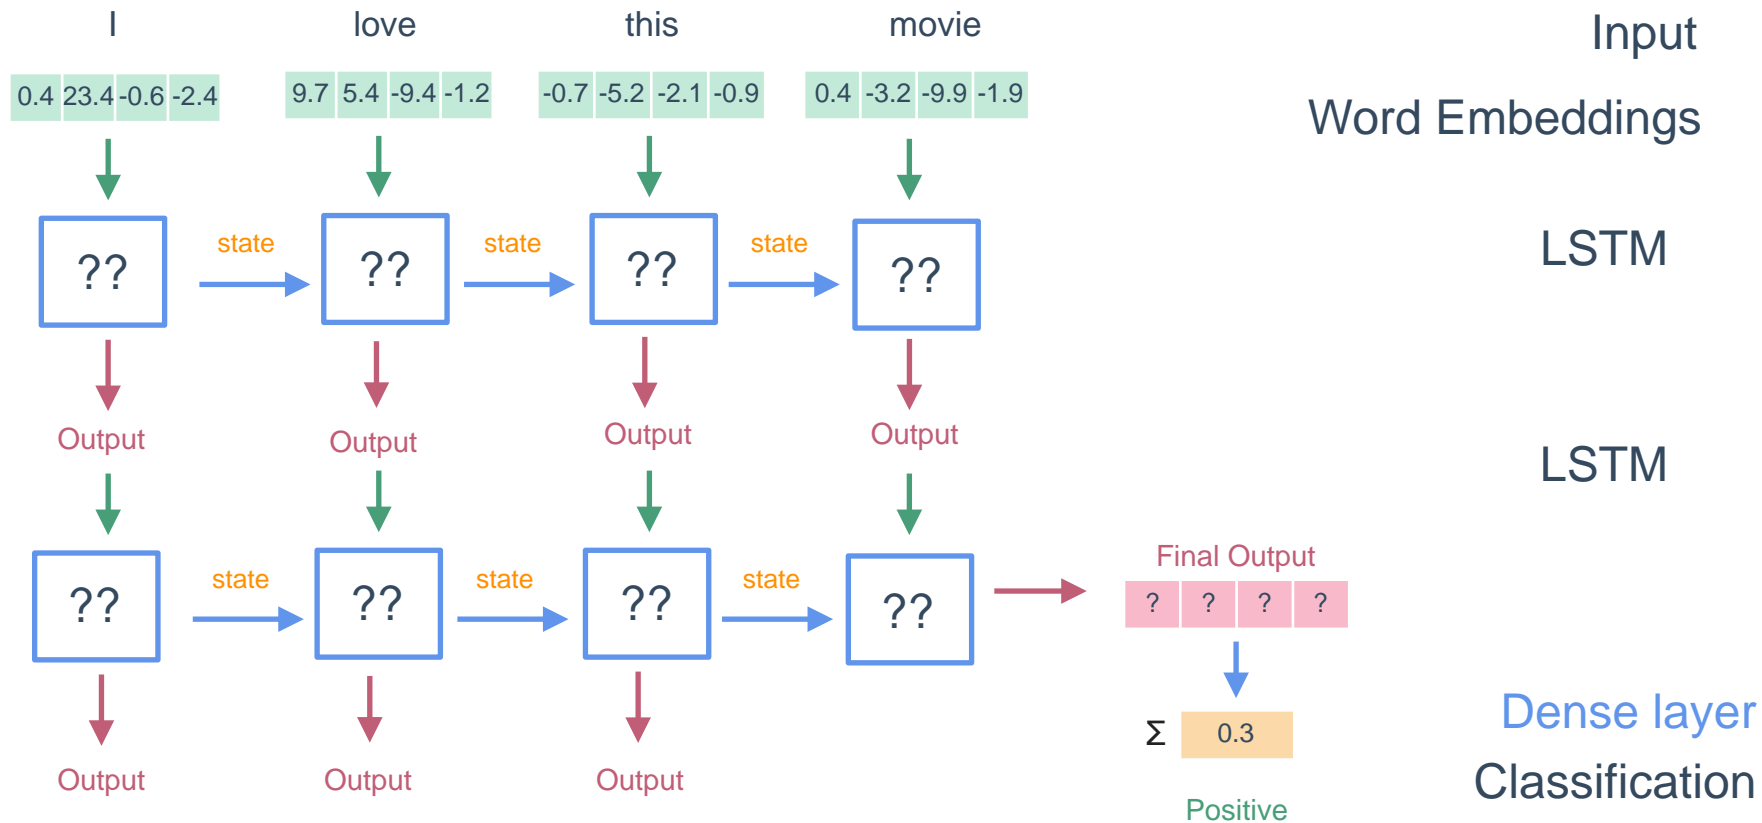
Deep LSTM



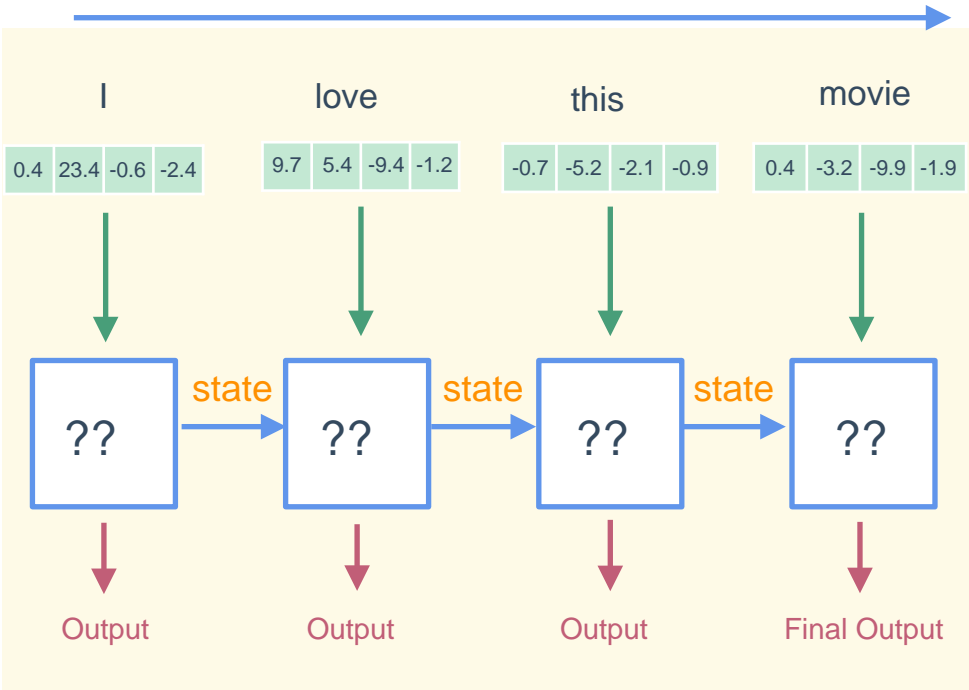
Input
Word Embeddings

LSTM

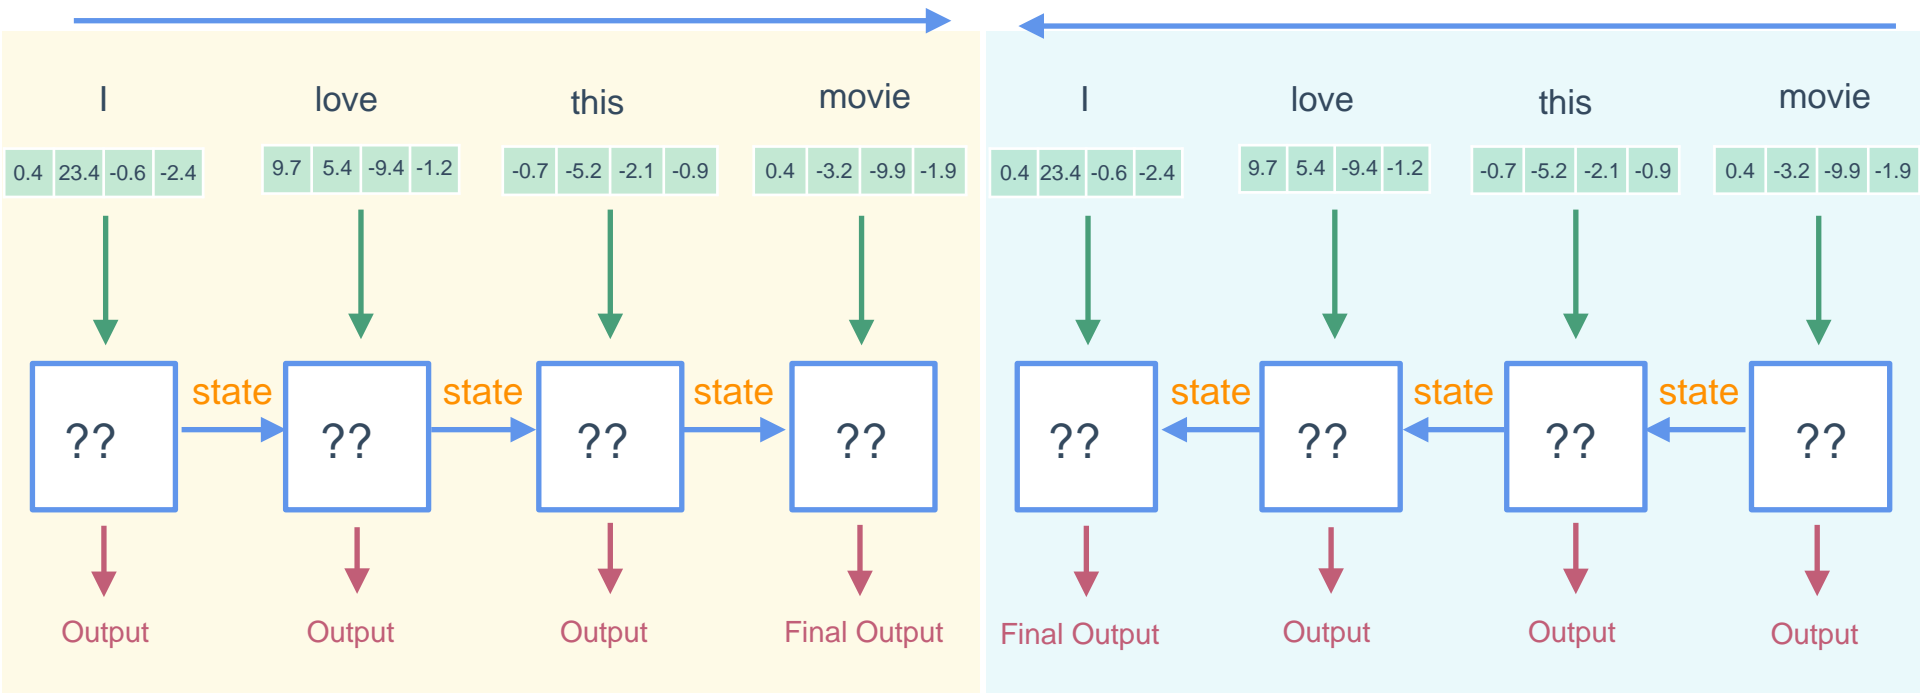
Deep LSTM



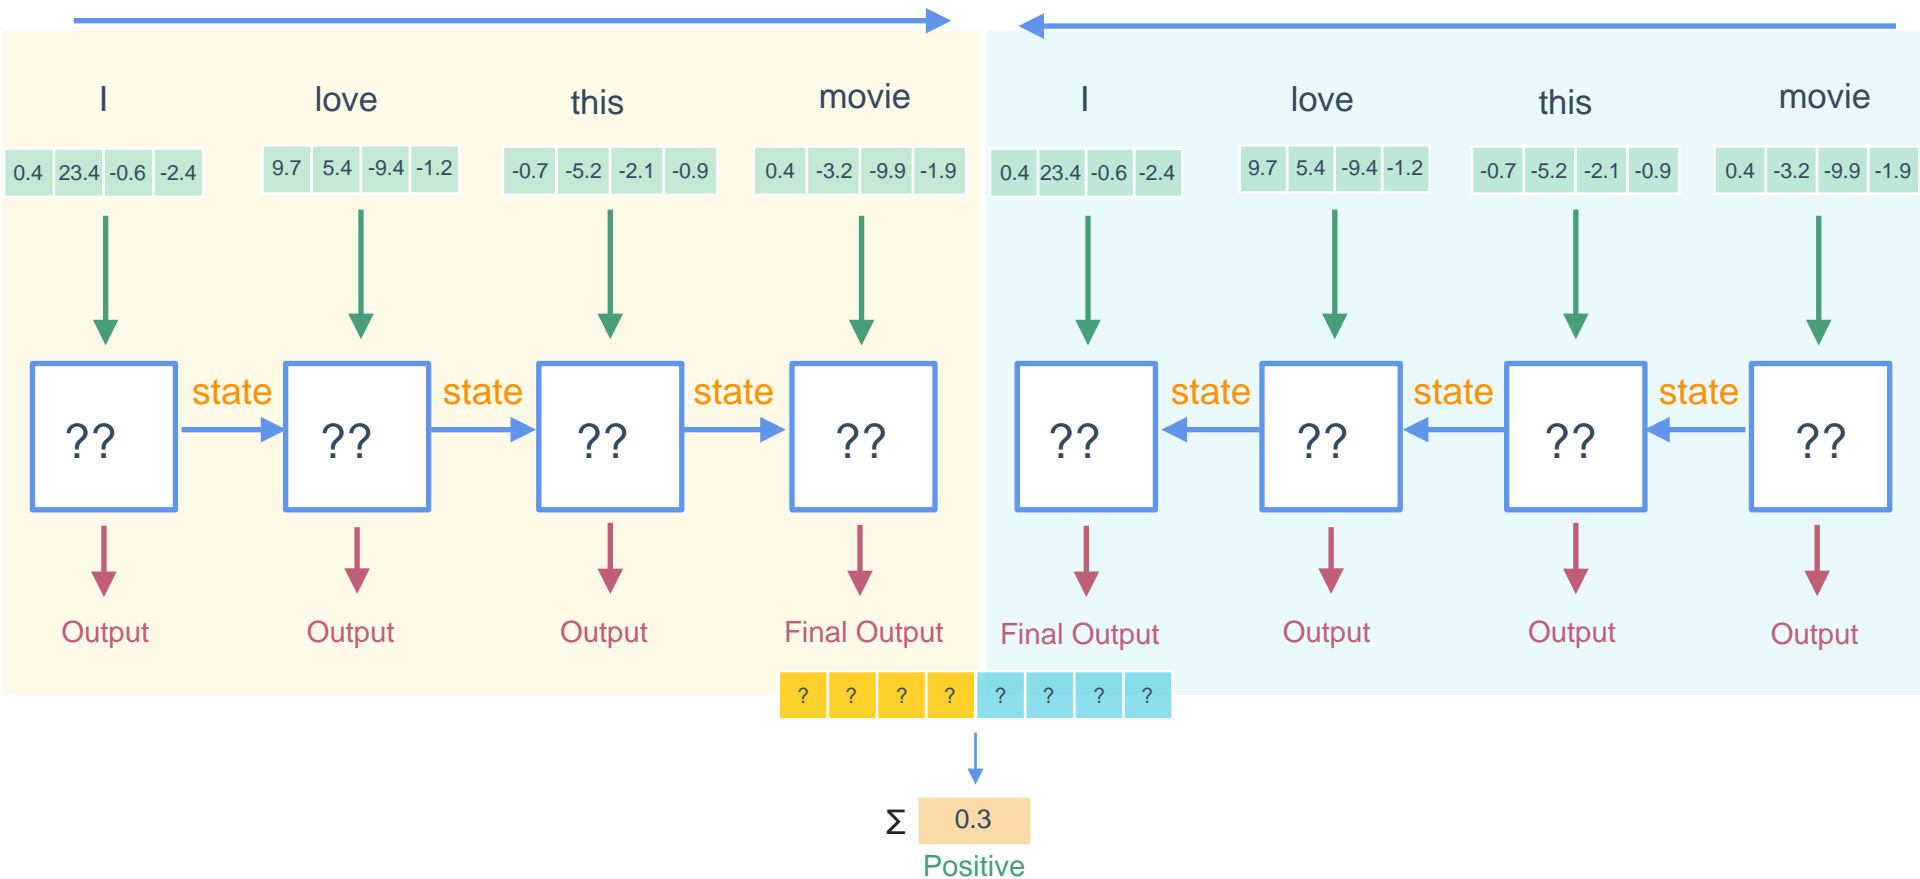
Bidirectional LSTM



Bidirectional LSTM



Bidirectional LSTM



The background is a solid blue color. Overlaid on this background is a faint, light-blue network diagram. This diagram consists of numerous small circular nodes connected by thin, straight lines. The nodes are arranged in a way that suggests a hierarchical or clustered structure, with some nodes acting as central hubs connected to many other nodes, while others are more isolated or part of smaller clusters. The overall effect is a technical, data-oriented aesthetic.

Text Classification with CNNs

2D Convolution Review

52	34	14	5
45 1	12 0	17 0.2	11
29 0.4	20 0.7	19 0.1	27
99 0.2	85 0	60 0.8	55
120	112	88	29

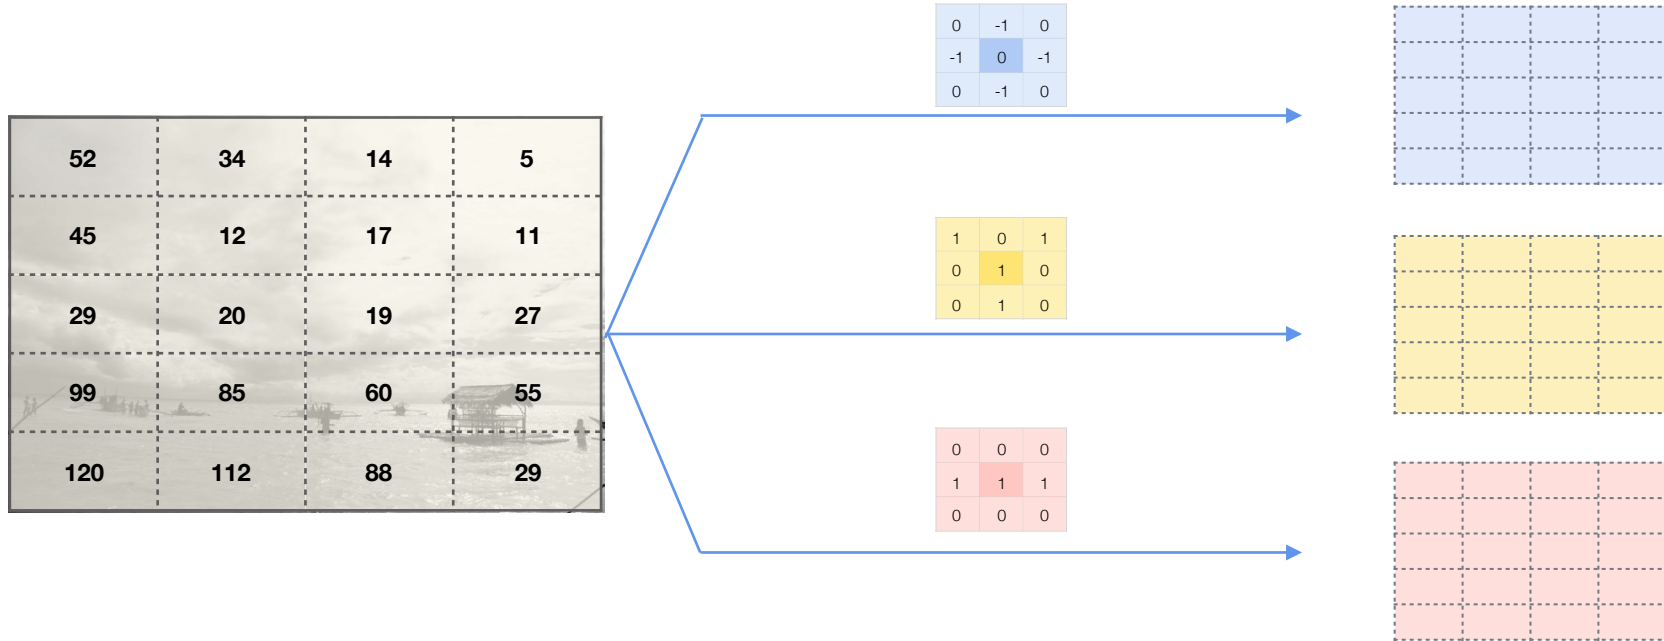
Input * Kernel

$$\begin{array}{r} (45 \times 1) \\ (12 \times 0) \\ (17 \times 0.2) \\ (29 \times 0.4) \\ (20 \times 0.7) \\ (19 \times 0.1) \\ (99 \times 0.2) \\ (85 \times 0) \\ (60 \times 0.8) \\ \hline + \\ \hline 143.7 \end{array}$$

131	122	97	17
114	104	78	88
	144		

Output

2D Convolution Review Multiple Outputs




Original Image

3 Convolutions with
different kernels

Multiple Outputs

2D Convolution Review Multiple Inputs



1D Convolution

	I		love		this		movie
Channel 1	0.4 0.2		9.7 0.3		-0.7 0.2		0.4
Channel 2	23.4		5.4		-5.2		-3.2
Channel 3	-0.6		-9.4		-2.1		-9.9
Channel 4	-2.4		-1.2		-0.9		-1.9

$$0.4 * 0.2 + 9.7 * 0.3 + (-0.7) * 0.2 = 2.85$$

Conv 1
2.85

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4 0.2	5.4 0.3	-5.2 0.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

$$23.4 * 0.2 + 5.4 * 0.3 + (-5.2) * 0.2 = 5.26$$

Conv 1
2.85
5.26

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6 <small>0.2</small>		-9.4 <small>0.3</small>	-2.1 <small>0.2</small>
Channel 4	-2.4	-1.2	-0.9	-1.9

$$(-0.6) \cdot 0.2 + (-9.4) \cdot 0.3 + (-2.1) \cdot 0.2 = -3.36$$

Conv 1
2.85
5.26
-3.36

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4 <small>0.2</small>		-1.2 <small>0.3</small>	-0.9 <small>0.2</small>

$$(-2.4) \cdot 0.2 + (-1.2) \cdot 0.3 + (-0.9) \cdot 0.2 = -1.02$$

Conv 1
2.85
5.26
-3.36
-1.02

1D Convolution

	I	love		this		movie	
Channel 1	0.4	9.7 0.3		-0.7 0.3		0.4 0.1	
Channel 2	23.4	5.4		-5.2		-3.2	
Channel 3	-0.6	-9.4		-2.1		-9.9	
Channel 4	-2.4	-1.2		-0.9		-1.9	
		Conv 1			Conv 2		
		2.85			2.74		
		5.26					
		-3.36					
		-1.02					

$$9.7 * 0.3 + (-0.7) * 0.3 + 0.4 * 0.1 = 2.74$$

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4 0.3	-5.2 0.3	-3.2 0.1
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9
		Conv 1	Conv 2	
		2.85	2.74	
		5.26	-0.26	
		-3.36		
		-1.02		

$$5.4 * 0.3 + (-5.2) * 0.3 + (-3.2) * 0.1 = -0.26$$

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4 <small>0.3</small>	-2.1 <small>0.3</small>	-9.9 <small>0.1</small>
Channel 4	-2.4	-1.2	-0.9	-1.9

Conv 1	Conv 2
2.85	2.74
5.26	-0.26
-3.36	-4.44
-1.02	

$$(-9.4) \cdot 0.3 + (-2.1) \cdot 0.3 + (-9.9) \cdot 0.1 = -4.44$$

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2 0.3	-0.9 0.3	-1.9 0.1
		Conv 1	Conv 2	$(-1.2)*0.3 + (-0.9)*0.3 + (-1.9)*0.1 = -0.82$
		2.85	2.74	
		5.26	-0.26	
		-3.36	-4.44	
		-1.02	-0.82	

2D Max Pooling Review

52	34	14	5
45	12	17	11
29	20	19	27
99	85	60	55
118	103	180	192
120	112	88	29

Input 6x4

Output 3x2

52	17		

2D Max Pooling Review

52	34	14	5
45	12	17	11
29	20	19	27
99	85	60	55
118	103	180	192
120	112	88	29

Input 6x4

Output 3x2

52	17		
99			

2D Max Pooling Review

52	34	14	5
45	12	17	11
29	20	19	27
99	85	60	55
118	103	180	192
120	112	88	29

Input 6x4

Output 3x2

52	17		
99	65		

1D Max Pooling

	I	love	this	movie
Channel 1	0.4		-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1

9.7

Max Pooling 2

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1

9.7
23.4

Max Pooling 2

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6		-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1

9.7
23.4
-0.6

Max Pooling 2

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4		-0.9	-1.9

Max Pooling 1
9.7
23.4
-0.6
-1.2

Max Pooling 2

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1
9.7
23.4
-0.6
-1.2

Max Pooling 2
0.4

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1
9.7
23.4
-0.6
-1.2

Max Pooling 2
0.4
-3.2

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

Max Pooling 1
9.7
23.4
-0.6
-1.2

Max Pooling 2
0.4
-3.2
-2.1

1D Max Pooling

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2	-0.9	-1.9

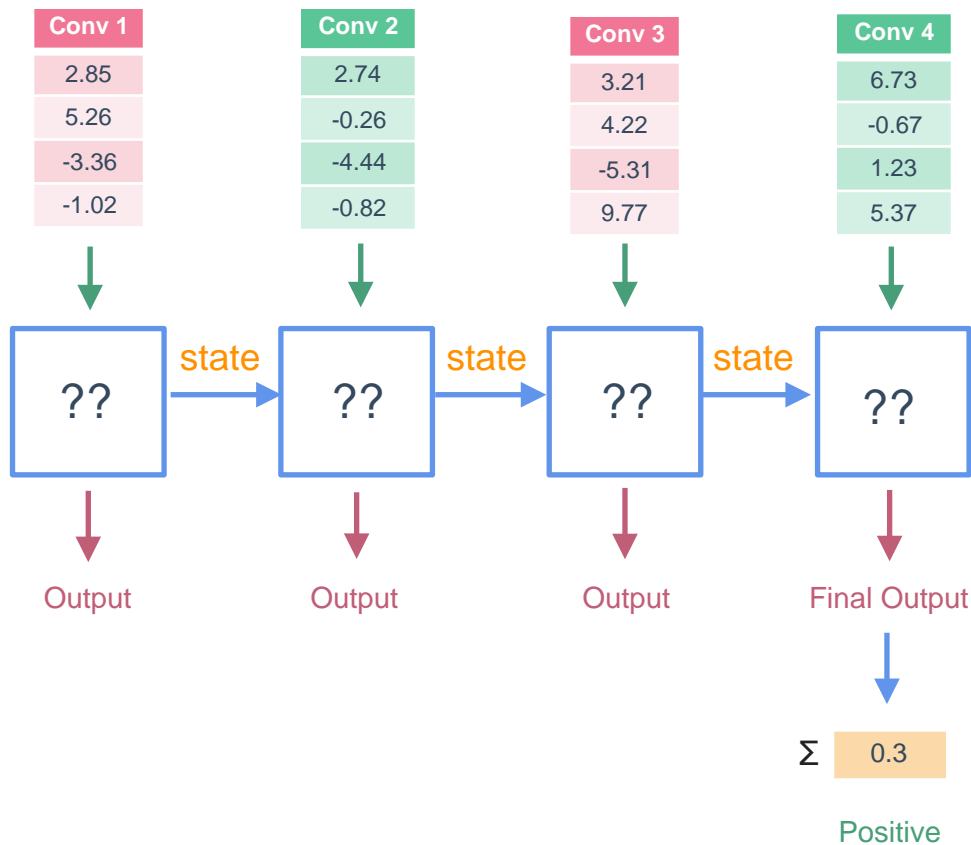
Max Pooling 1
9.7
23.4
-0.6
-1.2

Max Pooling 2
0.4
-3.2
-2.1
-0.9

1D Convolution

	I	love	this	movie
Channel 1	0.4	9.7	-0.7	0.4
Channel 2	23.4	5.4	-5.2	-3.2
Channel 3	-0.6	-9.4	-2.1	-9.9
Channel 4	-2.4	-1.2 0.3	-0.9 0.3	-1.9 0.1
		Conv 1	Conv 2	$(-1.2)*0.3 + (-0.9)*0.3 + (-1.9)*0.1 = -0.82$
		2.85	2.74	
		5.26	-0.26	
		-3.36	-4.44	
		-1.02	-0.82	

CNN/LSTM Hybrid



Input

Word embeddings


Convolutions

LSTM

Encoded Output

Dense layer

Classification



Seq2Seq
ml-class/lstm/seq2seq

Seq2Seq in translation

Input

My name is Lukas



??



私の名前はルーカスです

Output



seq2seq: the clown car of deep learning

tl; dr: Translating arbitrary-length sequences back and forth is easier than you think




Seq2Seq in this tutorial



Seq2Seq in this tutorial

Input

10 + 12



22

Output

1

0

+

1

2

Input

1

0

+

1

2

Input

0 0 1 0

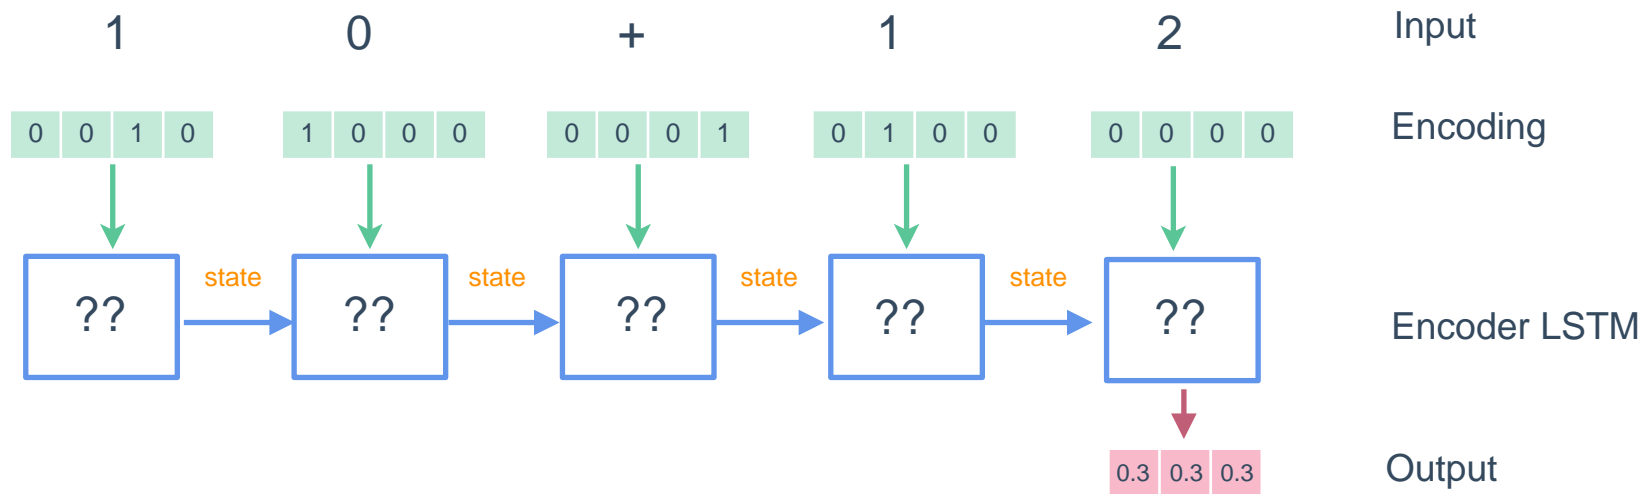
1 0 0 0

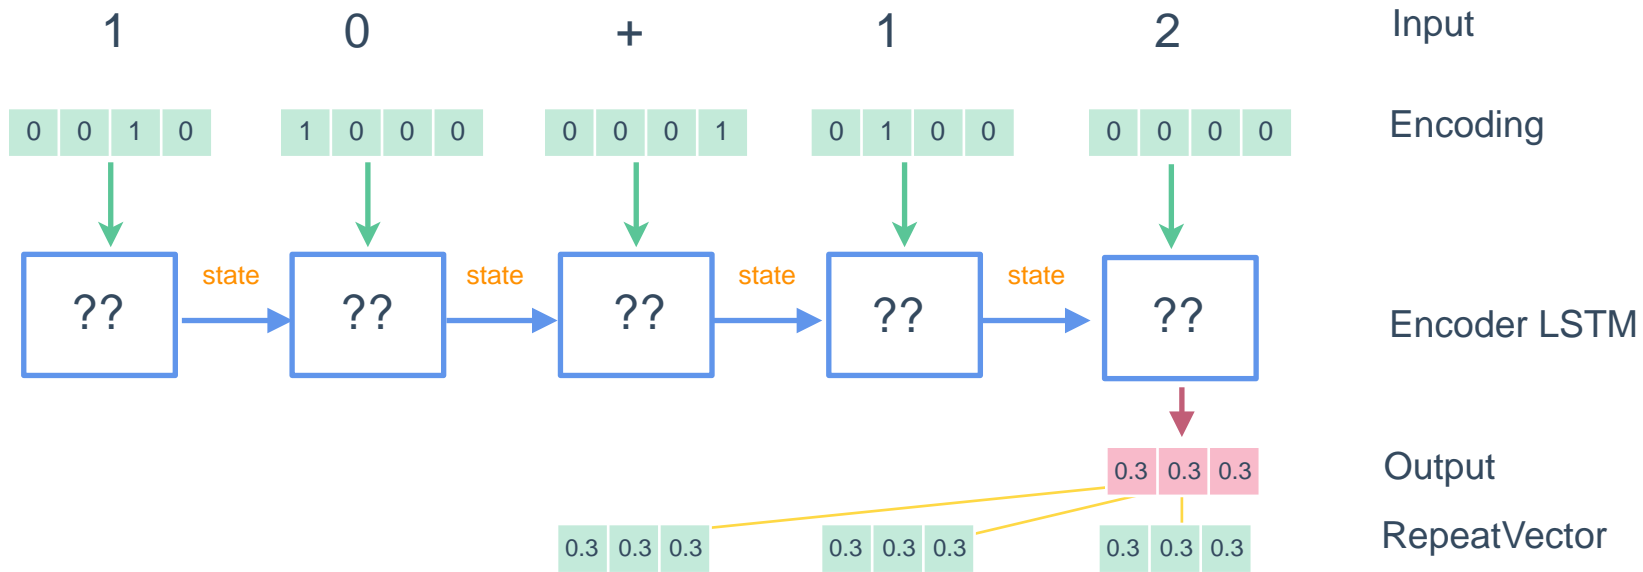
0 0 0 1

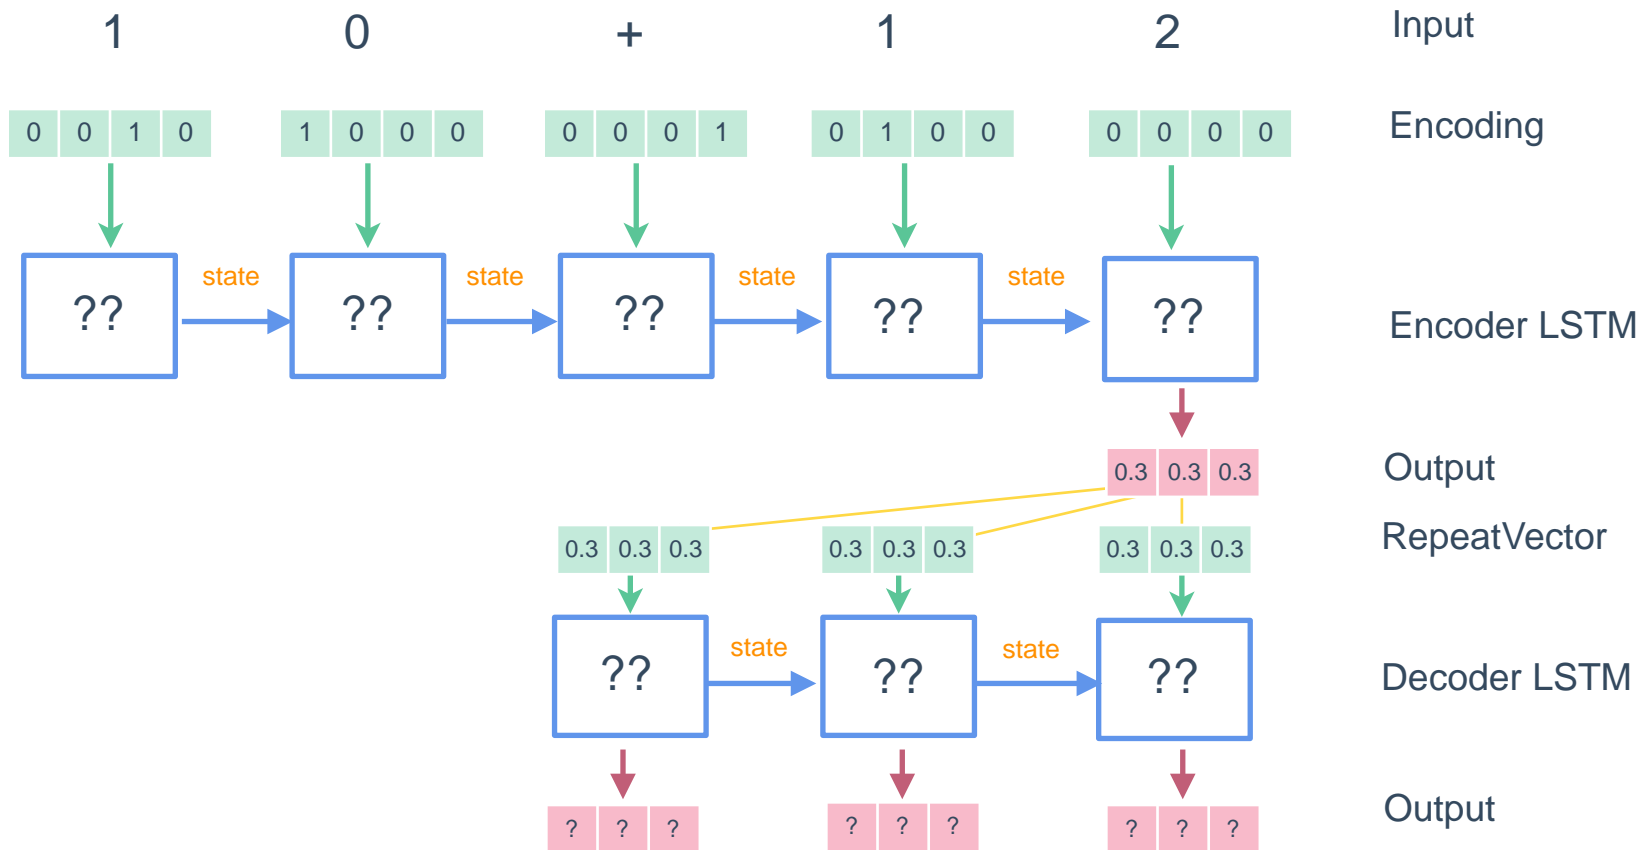
0 1 0 0

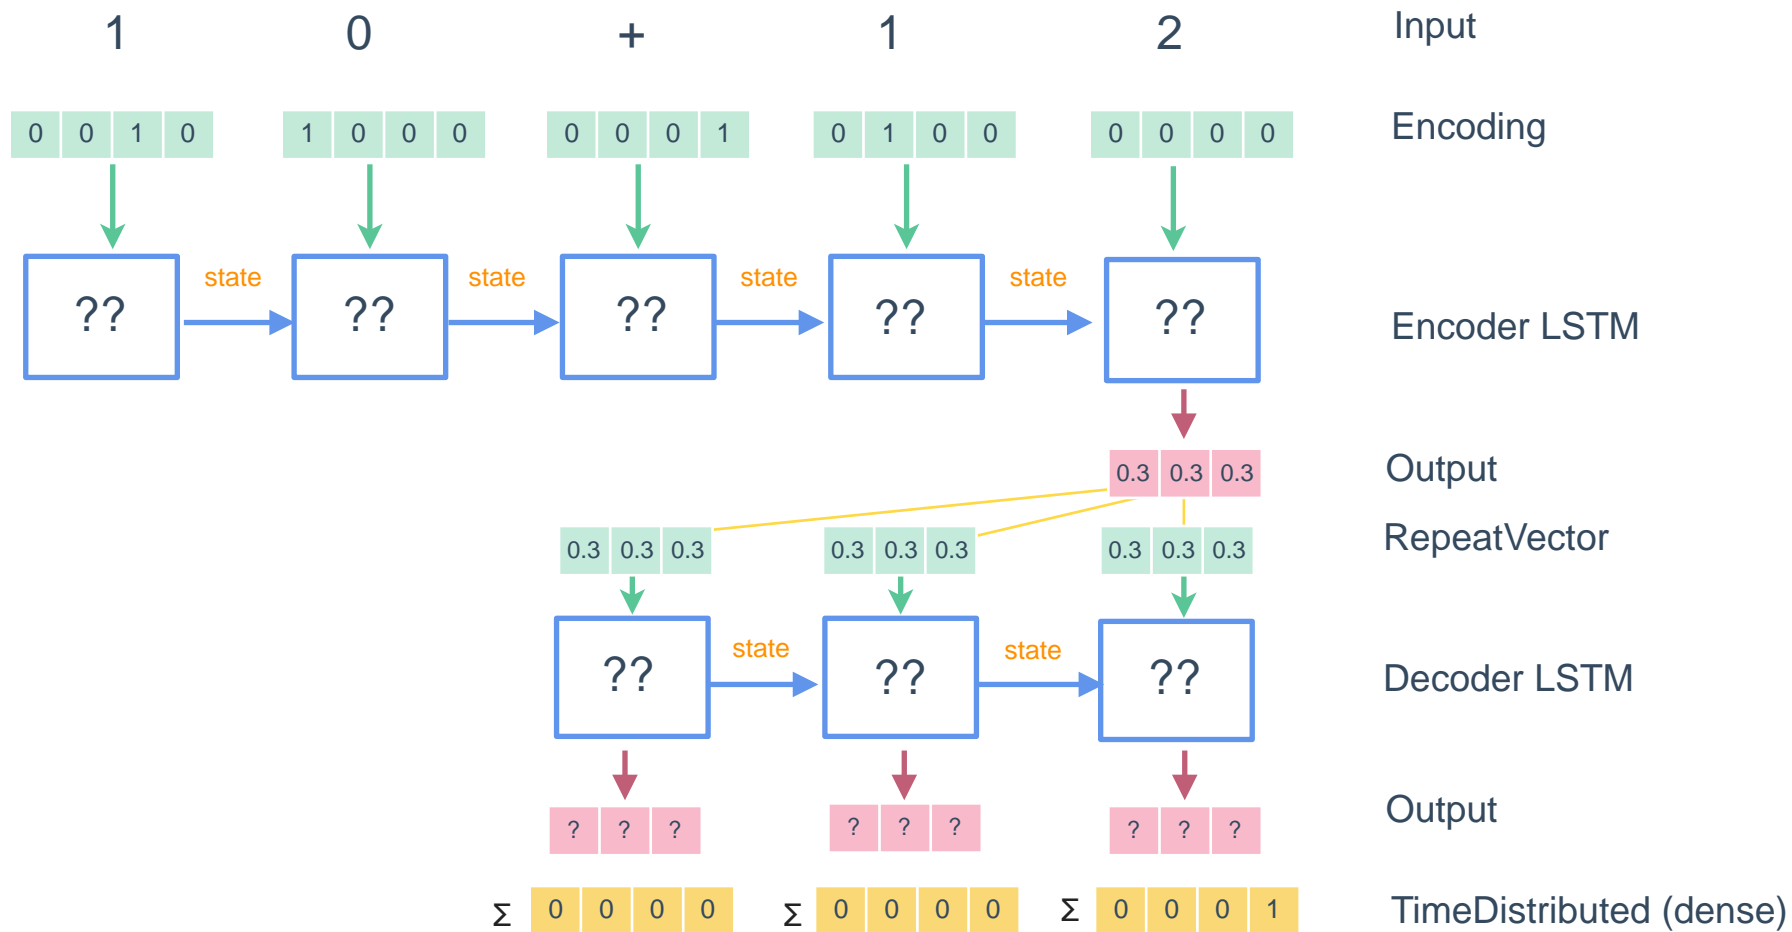
0 0 0 0

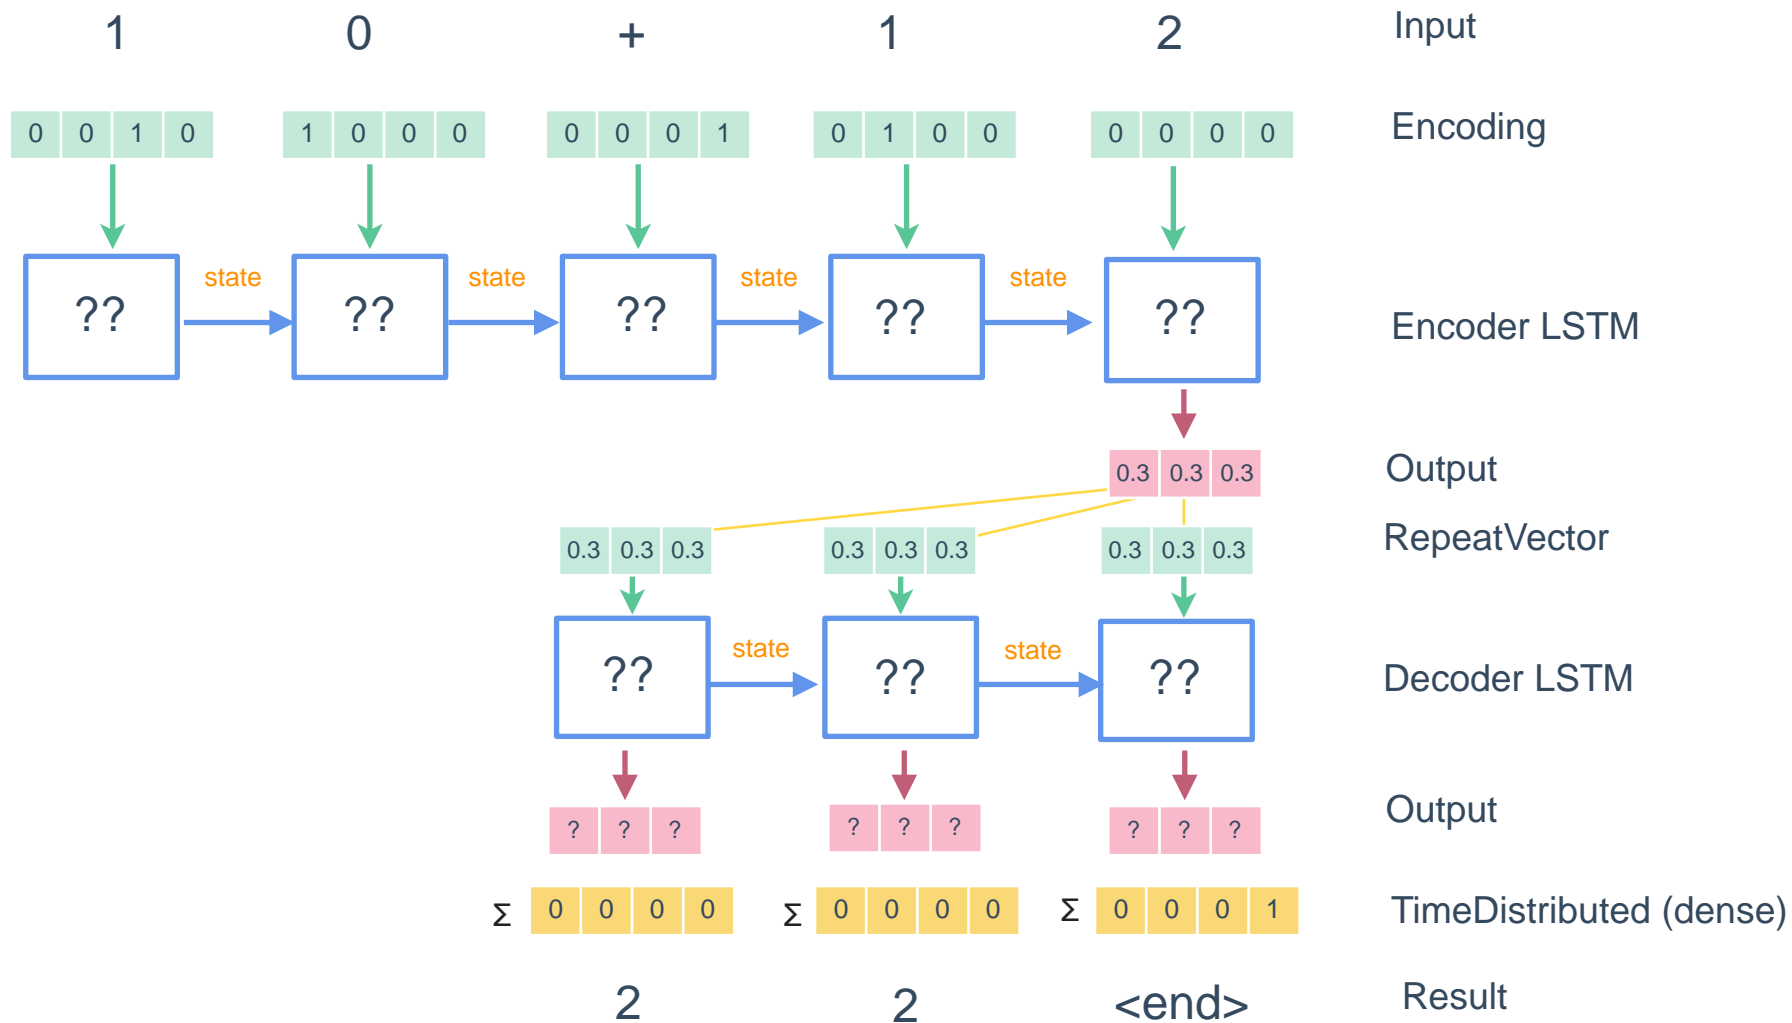
Encoding





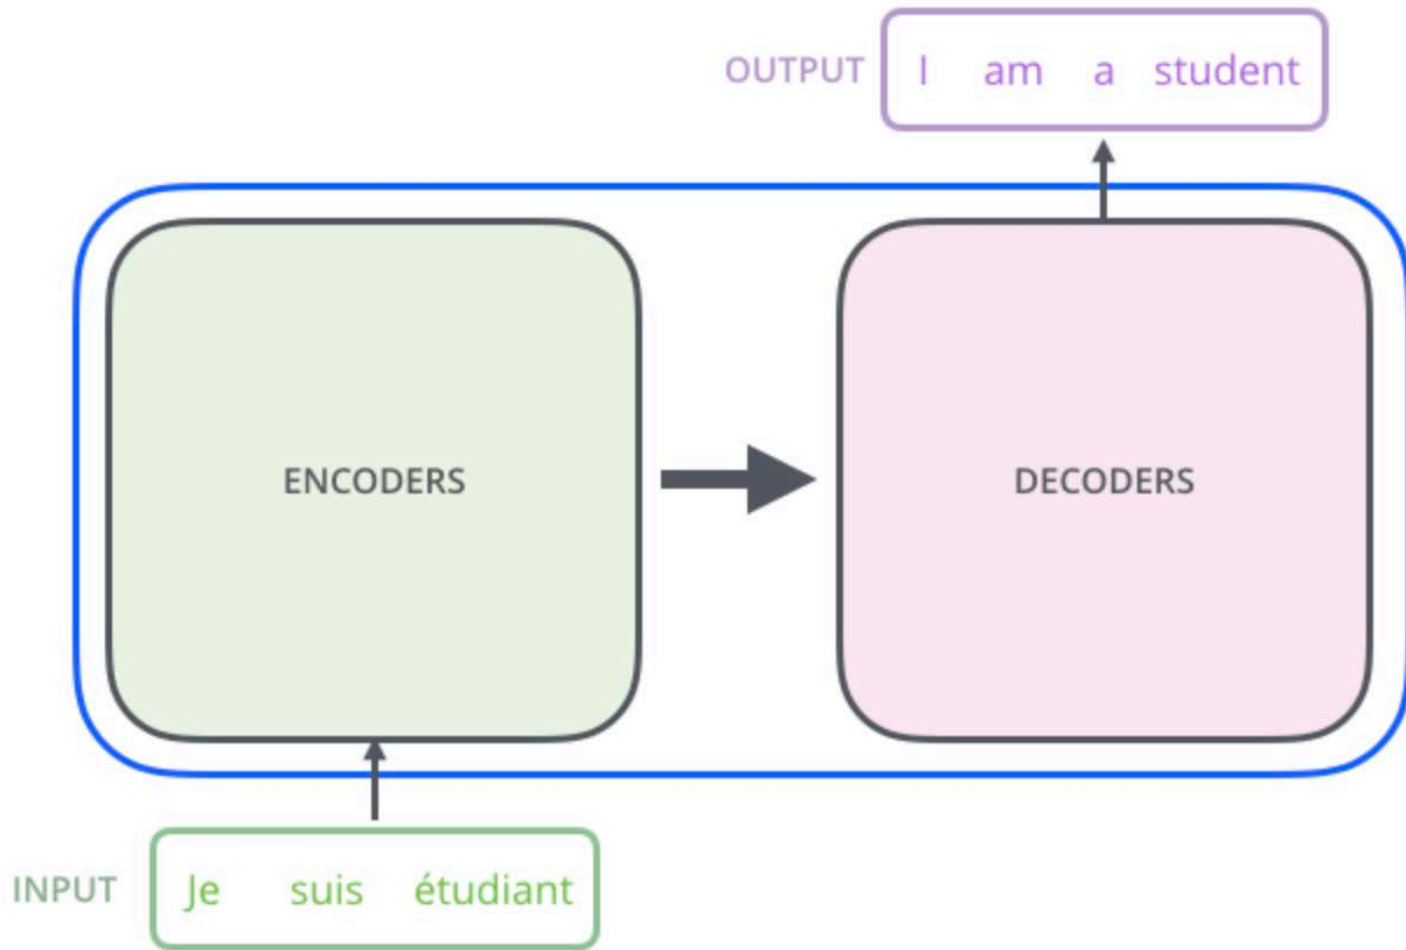




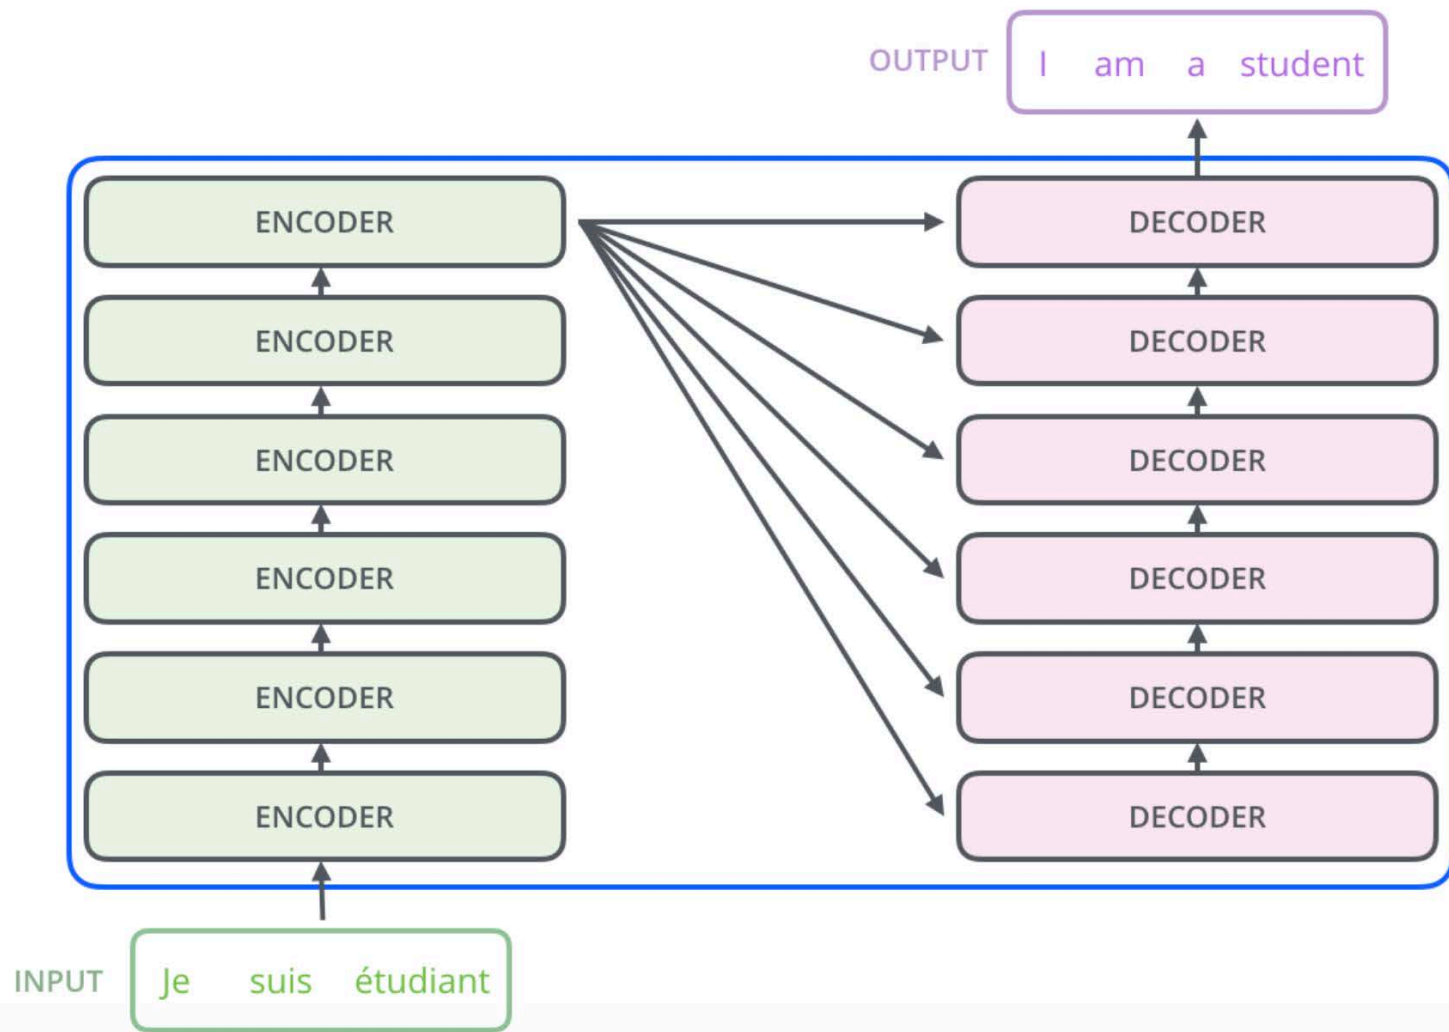


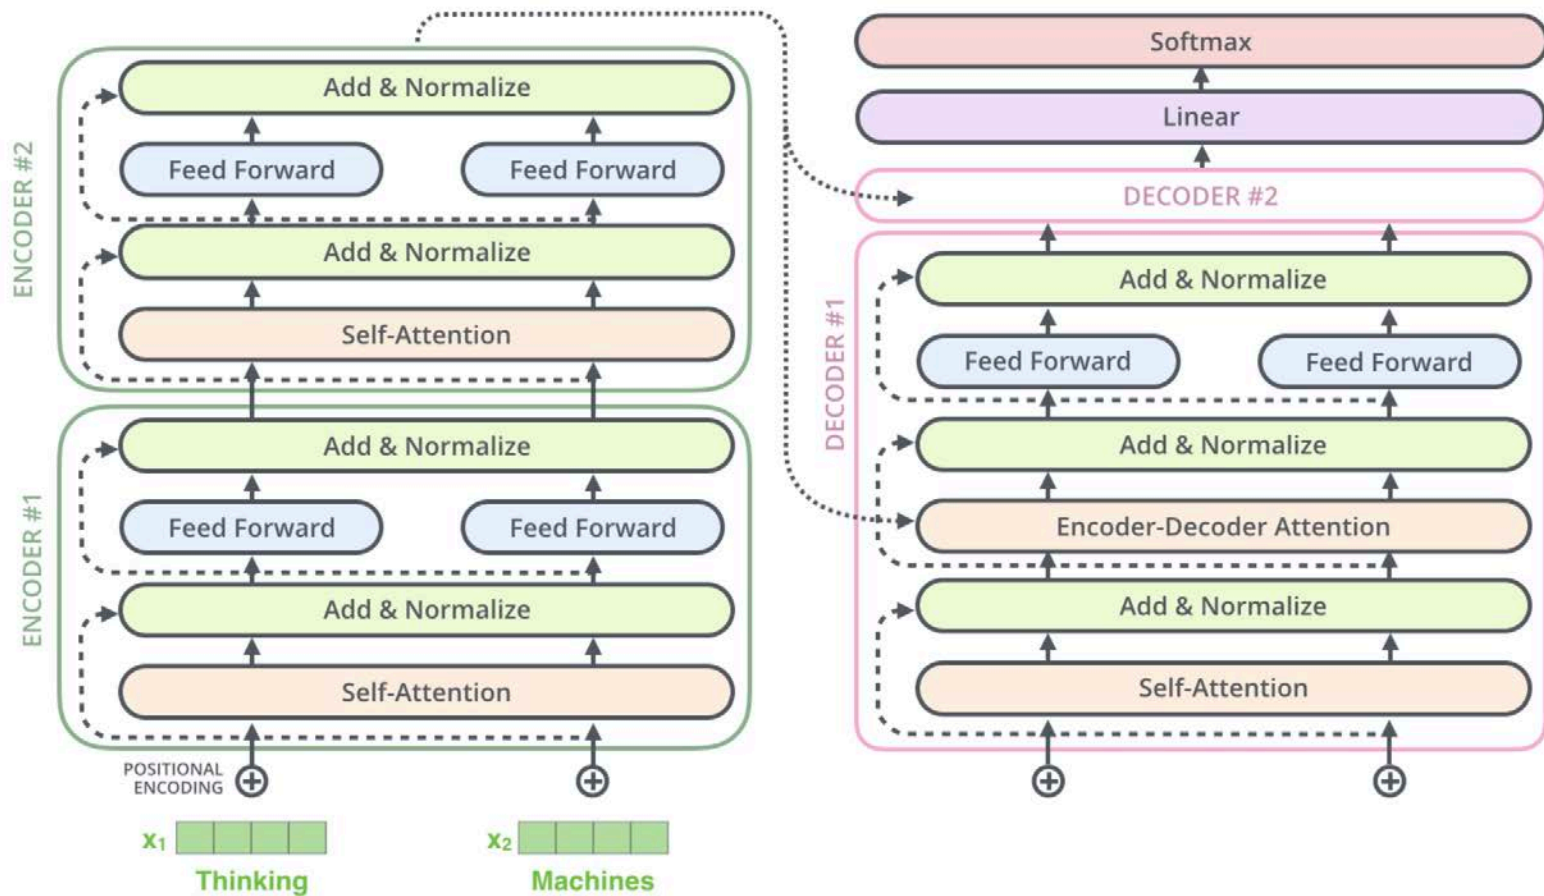
The background of the slide features a complex, abstract network diagram. It consists of numerous small, light-blue circular nodes. A large, dense cluster of these nodes is positioned in the center-left. Several other smaller, star-like clusters are distributed around the periphery. These clusters are interconnected by thin, light-blue lines, creating a web-like structure that suggests a global or interconnected network. The entire graphic is rendered in a monochromatic light blue color against a slightly darker blue background.

Transformers



<http://jalammar.github.io/illustrated-transformer/>





The background is a solid blue color. Overlaid on this is a faint, light-blue network diagram. It consists of numerous small circular nodes connected by thin, straight lines. The nodes are arranged in a way that suggests a hierarchical or radial structure, with some nodes acting as central hubs connected to many other nodes, while others are more isolated or part of smaller clusters. The overall effect is a complex, interconnected web of points and lines.

Conclusion

More about LSTMs

- More online tutorials <https://www.wandb.com/classes>
- Colah's blog <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- Andrej's blog <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- Attention networks <https://richliao.github.io/supervised/classification/2016/12/26/textclassifier-HATN/>
- Stanford CS 224N

More Resources for Deep Learning/ML

- Books
 - Deep Learning Book (<http://www.deeplearningbook.org/>)
 - Artificial Intelligence: A Modern Approach
 - Hands-On Machine Learning with Scikit-Learn and TensorFlow
 - Deep Learning with Python
- Video Classes
 - wandb.com/classes
- Hands-on
 - wandb.com/benchmarks